



VLDB 2024 - 50th International Conference on Very Large Data Bases


HAIChart: Human and AI Paired Visualization System

Yupeng Xie¹, Yuyu Luo^{1,2*}, Guoliang Li³, Nan Tang^{1,2}

¹HKUST (GZ), ²HKUST, ³Tsinghua University


Introduction

- Existing data visualization tools fall into **two main categories**:



Human Cost 💰💰💰
Result Quality ★★

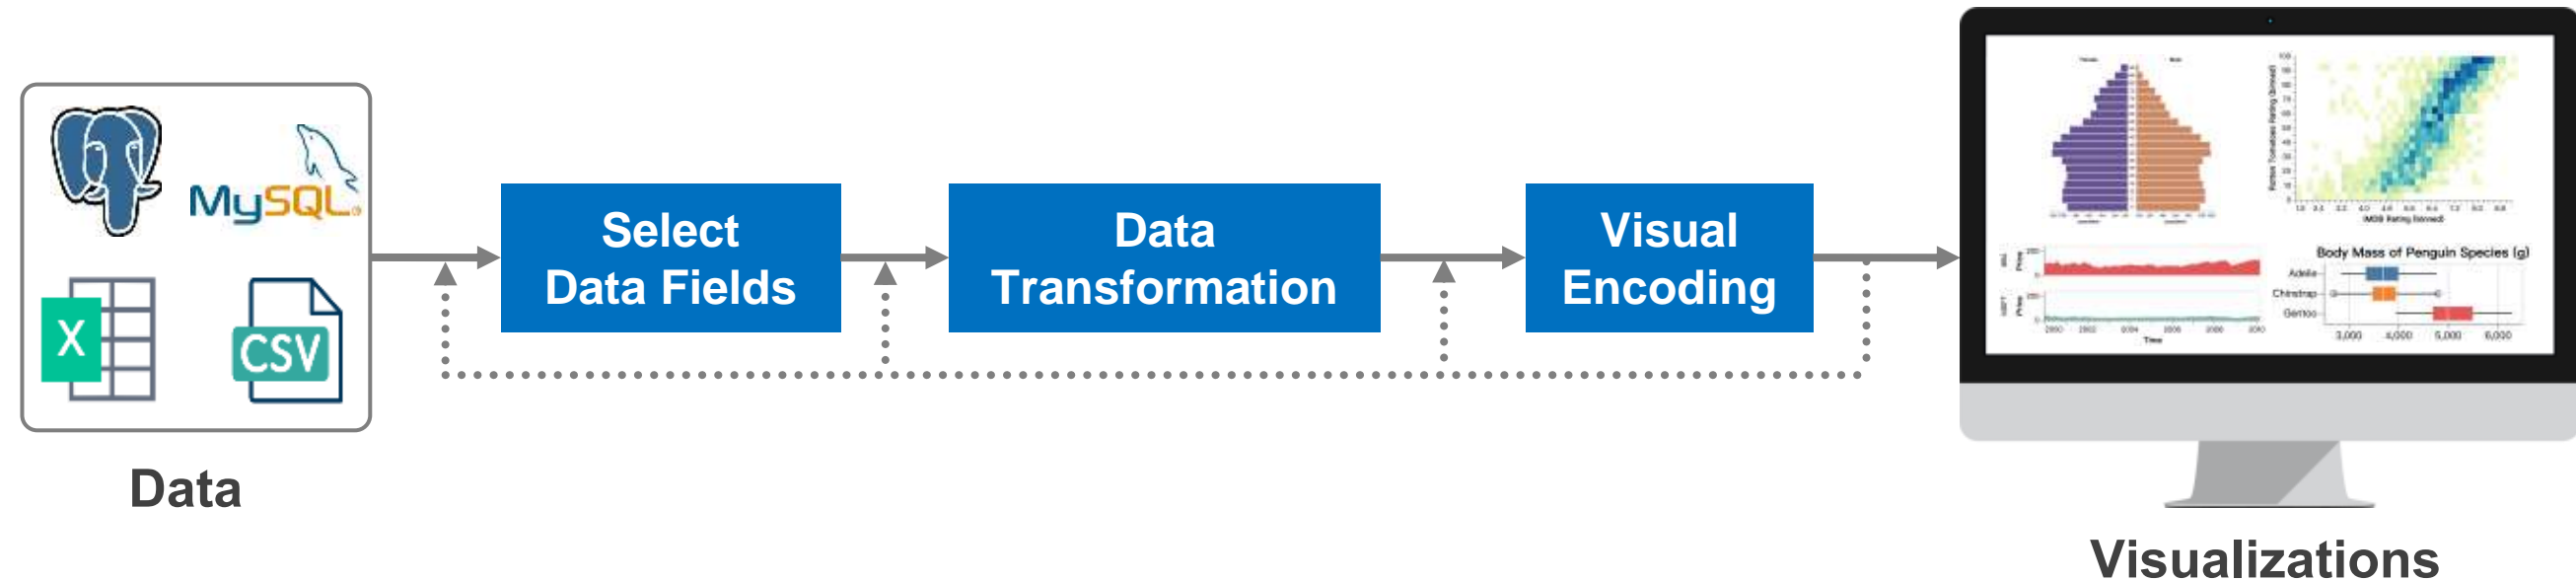
Human-powered VIS



Human Cost 💰
Result Quality ★★

Fully Automatic VIS

Human-powered Visualization



Manually specify through



Code


or



Clicks

- *Require human and domain expertise*
- *Tedious and **time-consuming** (even for experts)*


Fully Automatic Visualization



- Cannot capture **user intent or feedback**.
- Fails to meet specific **user needs**.

HAIChart: Human and AI Paired Visualization System

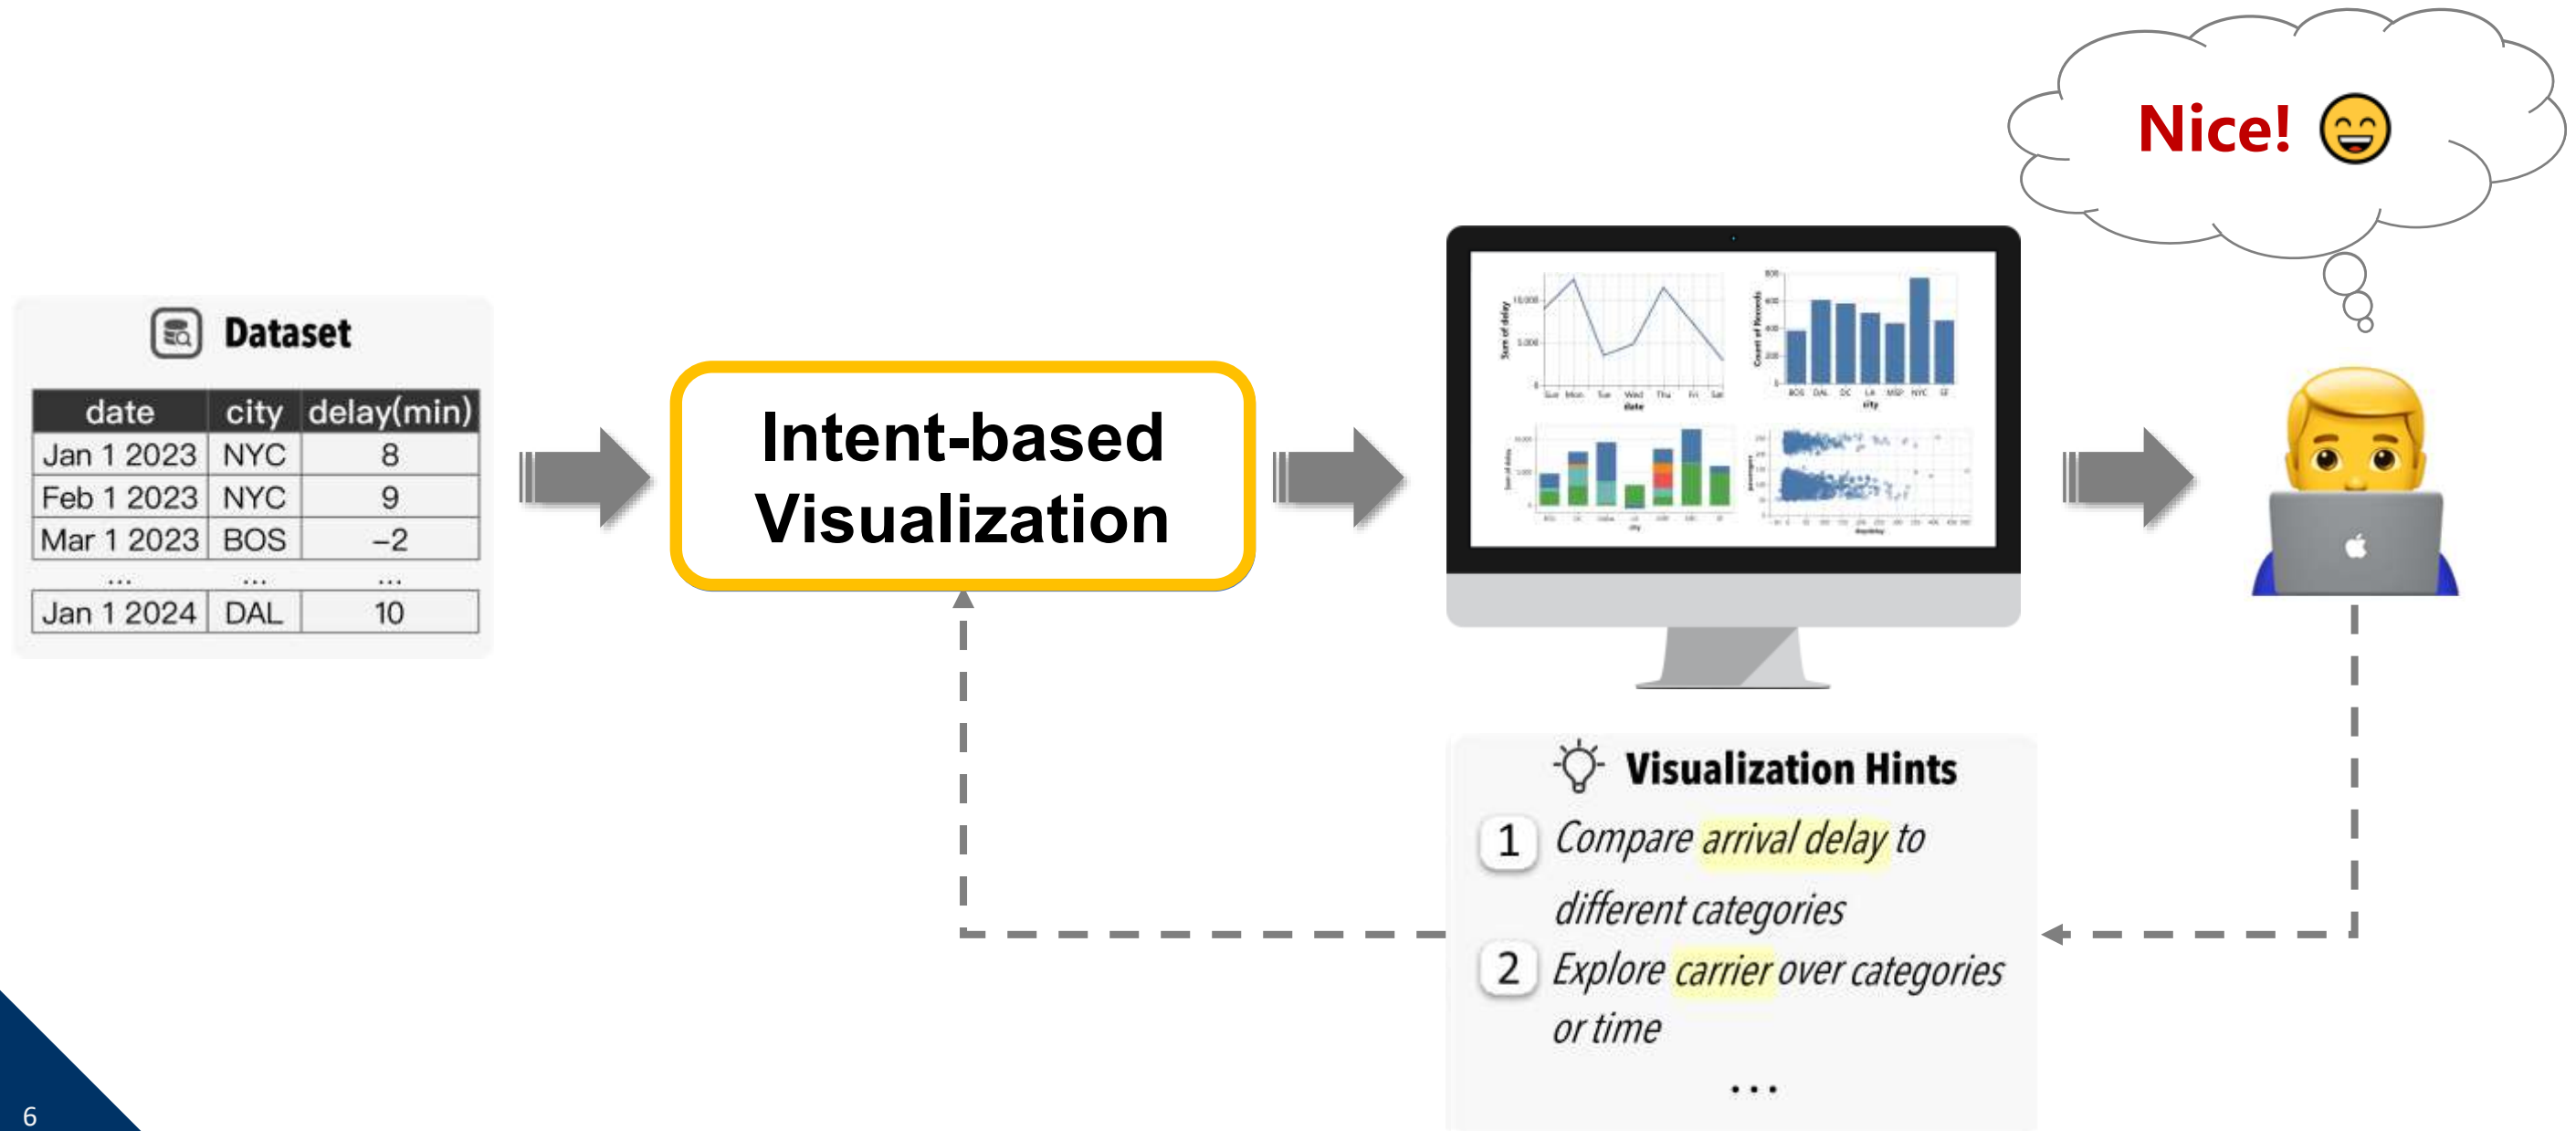
Phase 1: Recommend **high-quality** visualizations **to minimize manual effort.**



HAIChart: Human and AI Paired Visualization System

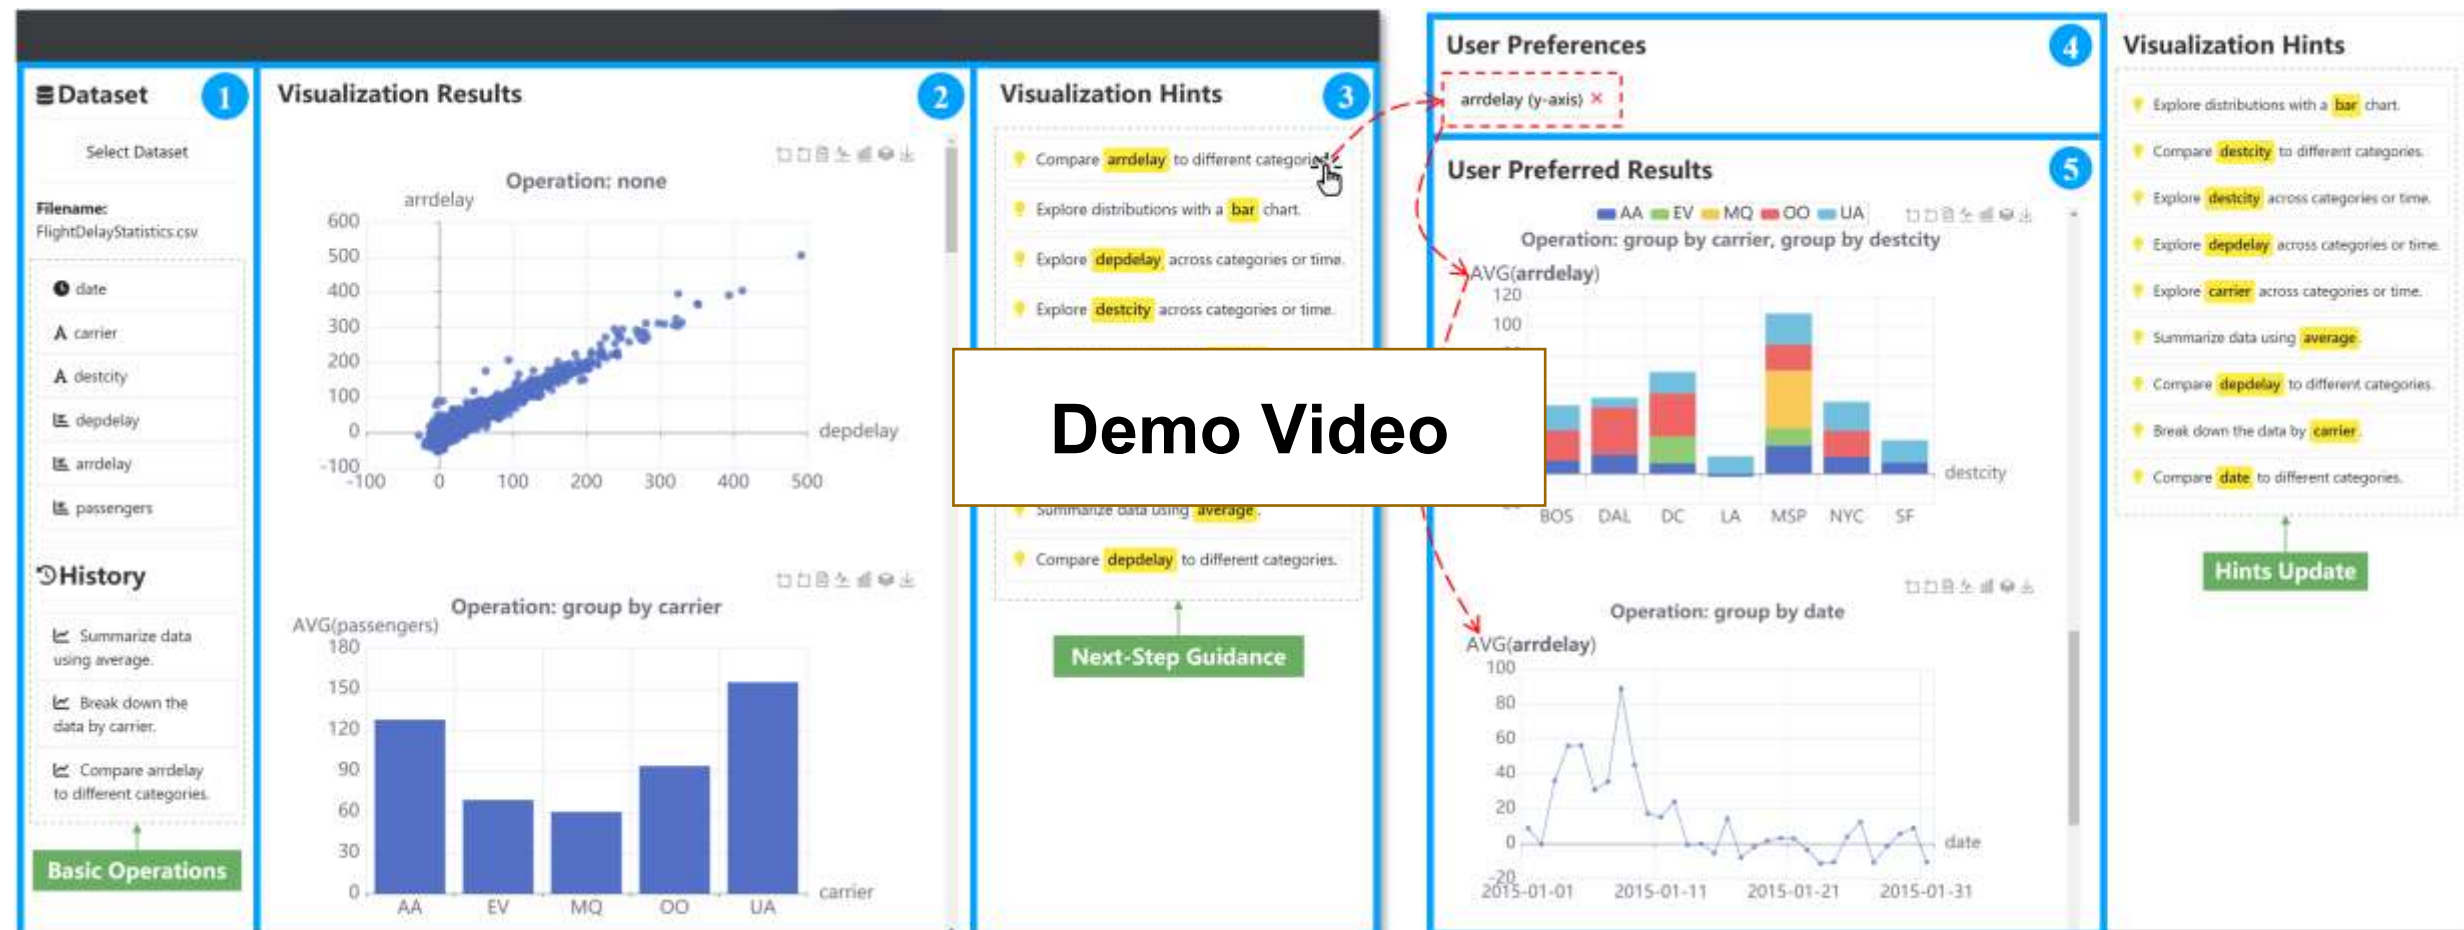
Phase 1: Recommend **high-quality** visualizations **to minimize manual effort.**

Phase 2: **Refine** visualizations **with hints** to more closely **align with user needs.**

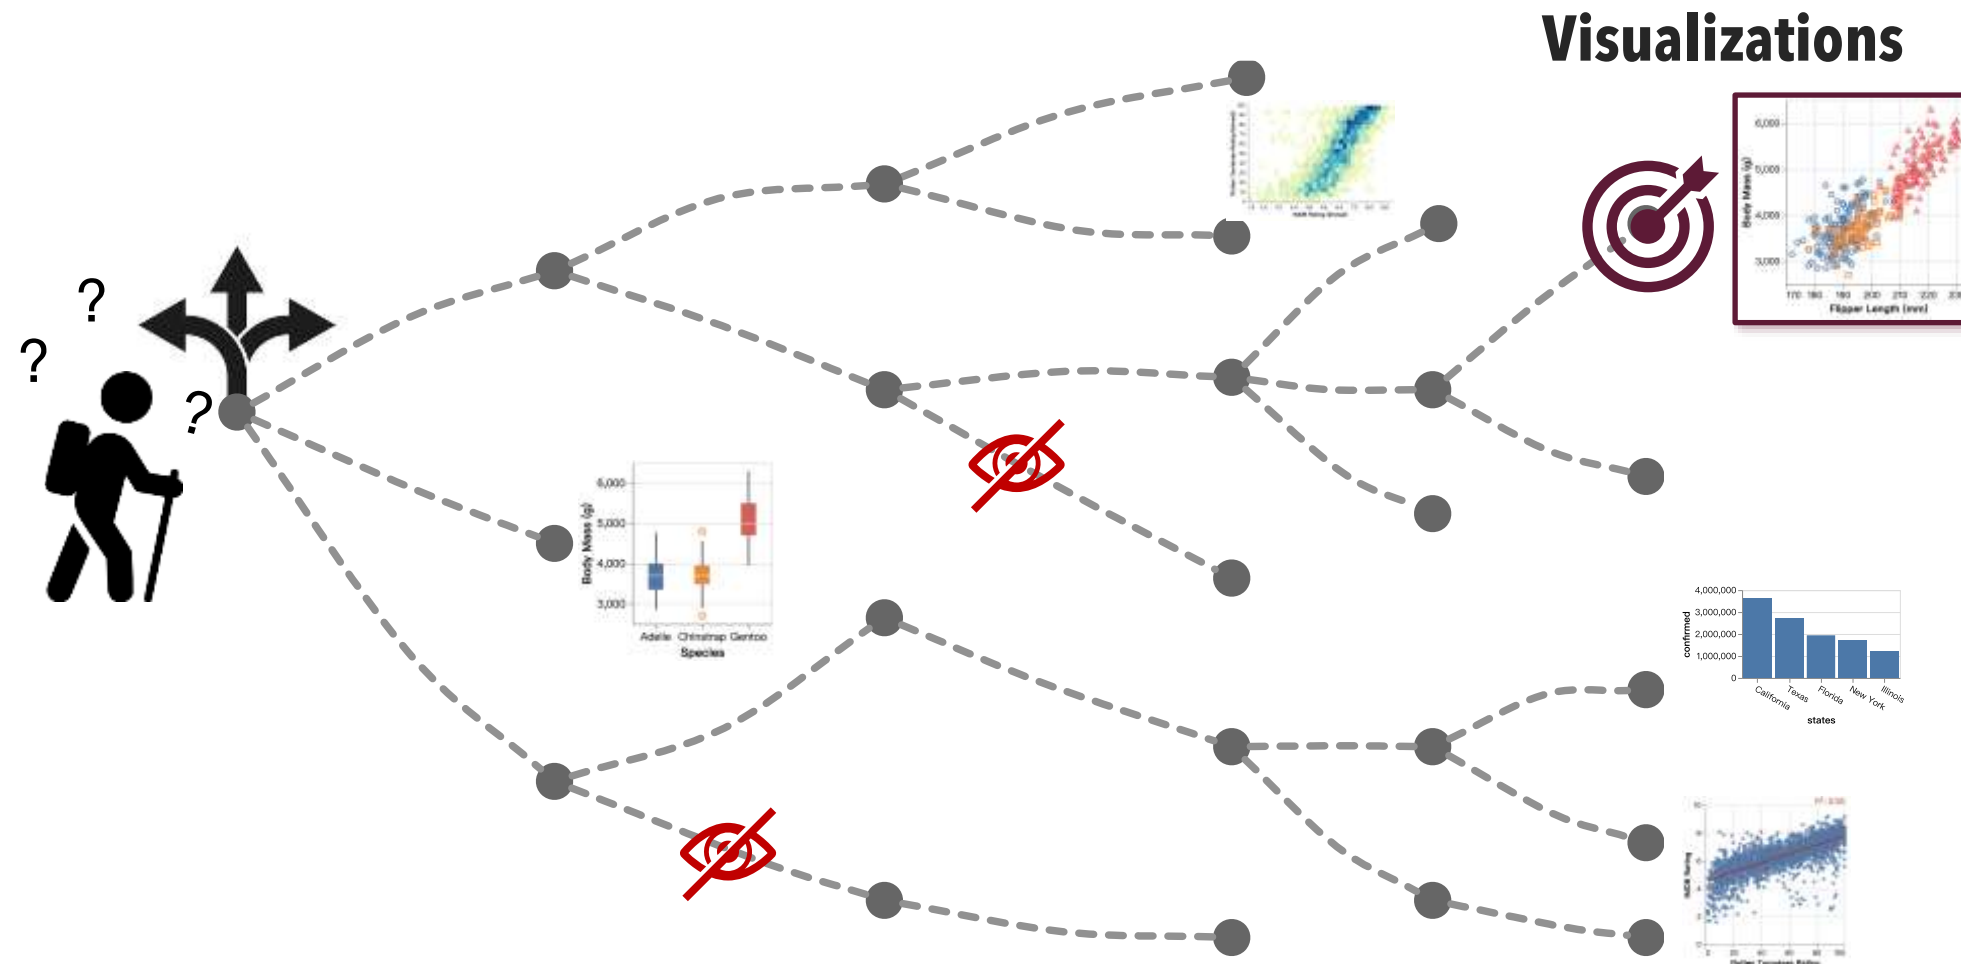


Demonstration Scenarios

- First-Round Visualization Recommendations.
- Multi-Round Visualization Recommendation based on Hints.



Challenge 1: How to Explore Search Space Efficiently?



Visualization Query

We use **visualization queries** to represent all possible visualizations.

Each query is a **sequence of operations** such as visual encoding and data transformation.

Flight Delay Dataset


Date	City	Delay(min)	...
Jan 1 2023	NYC	8	...
Feb 1 2023	NYC	9	...
Mar 1 2023	BOS	-2	...
...

mark
encoding
transform

Bar
x City y AVG(Delay)
group City

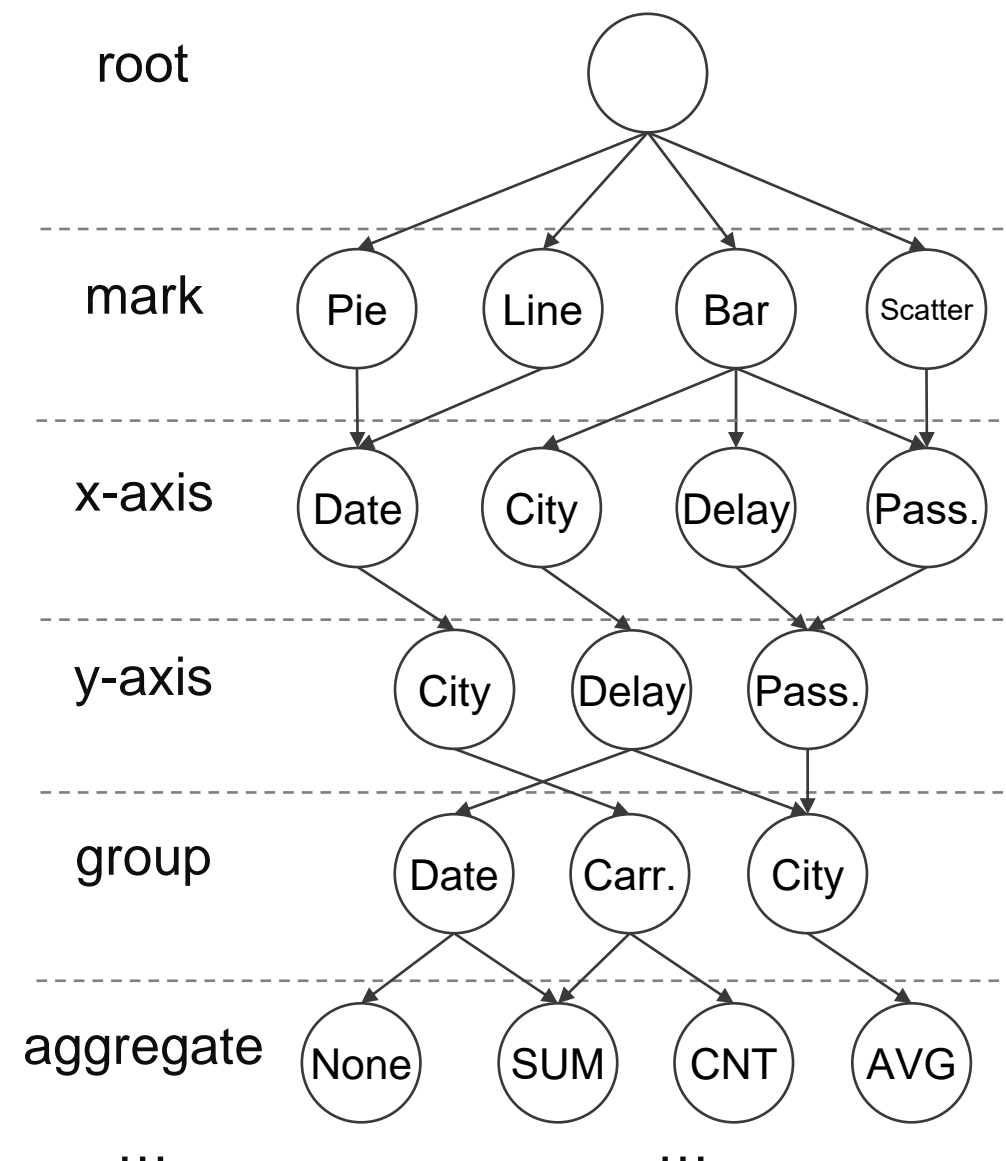
Query

Visualization



Visualization Query Graph


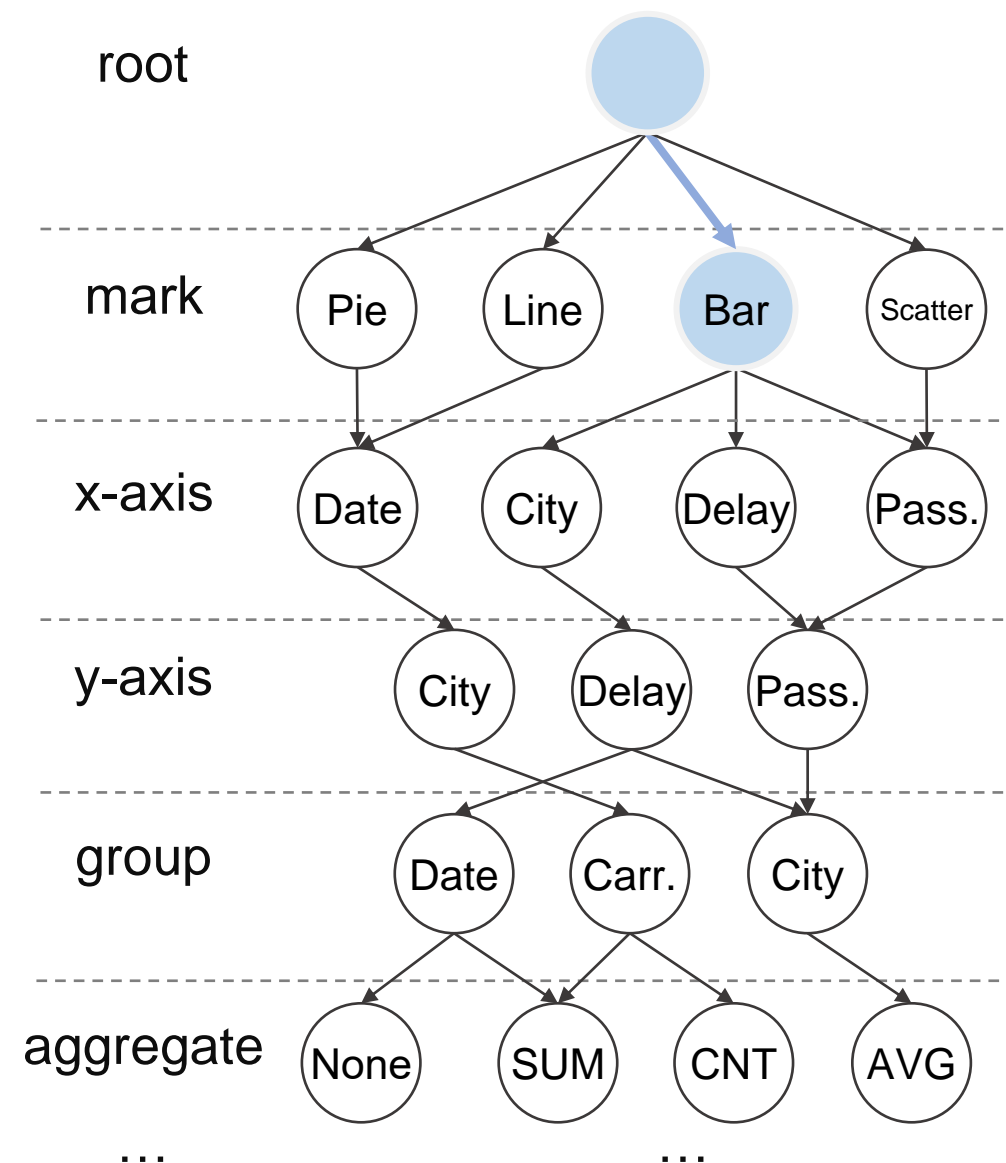
We also introduce the **visualization query graph** to explore various visualizations.



- Each **node** represents a visualization **operation**.
- A **path** is a sequence of these operations, representing a **query**.

Visualization Query Graph


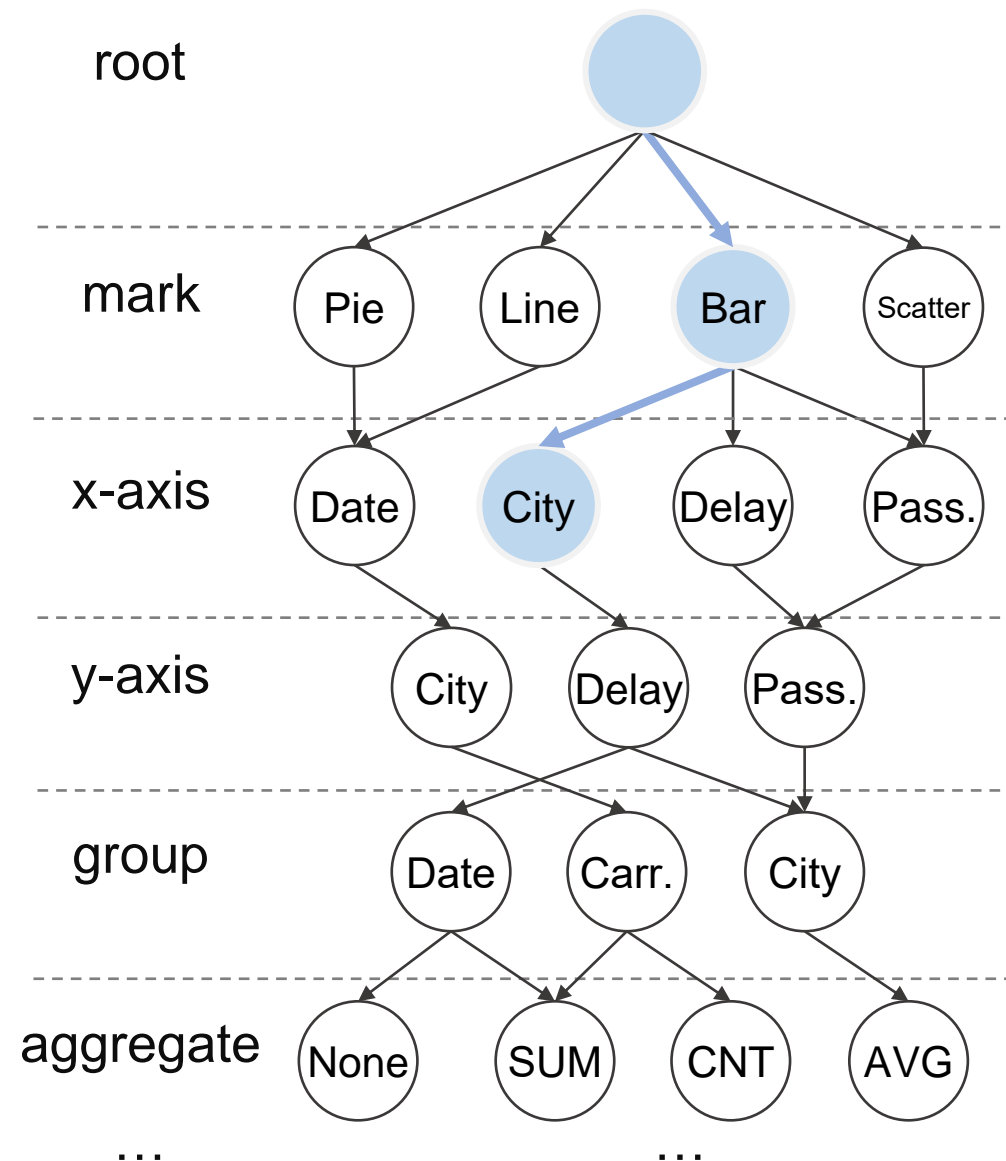
For example, **start from the root node**. First, select the chart type as a **bar chart**.



Current Visualization

Visualization Query Graph

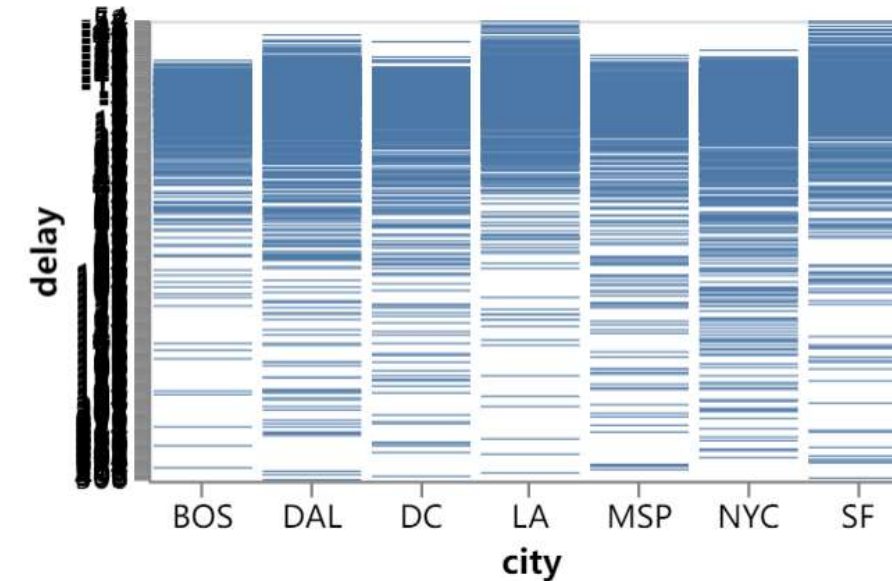
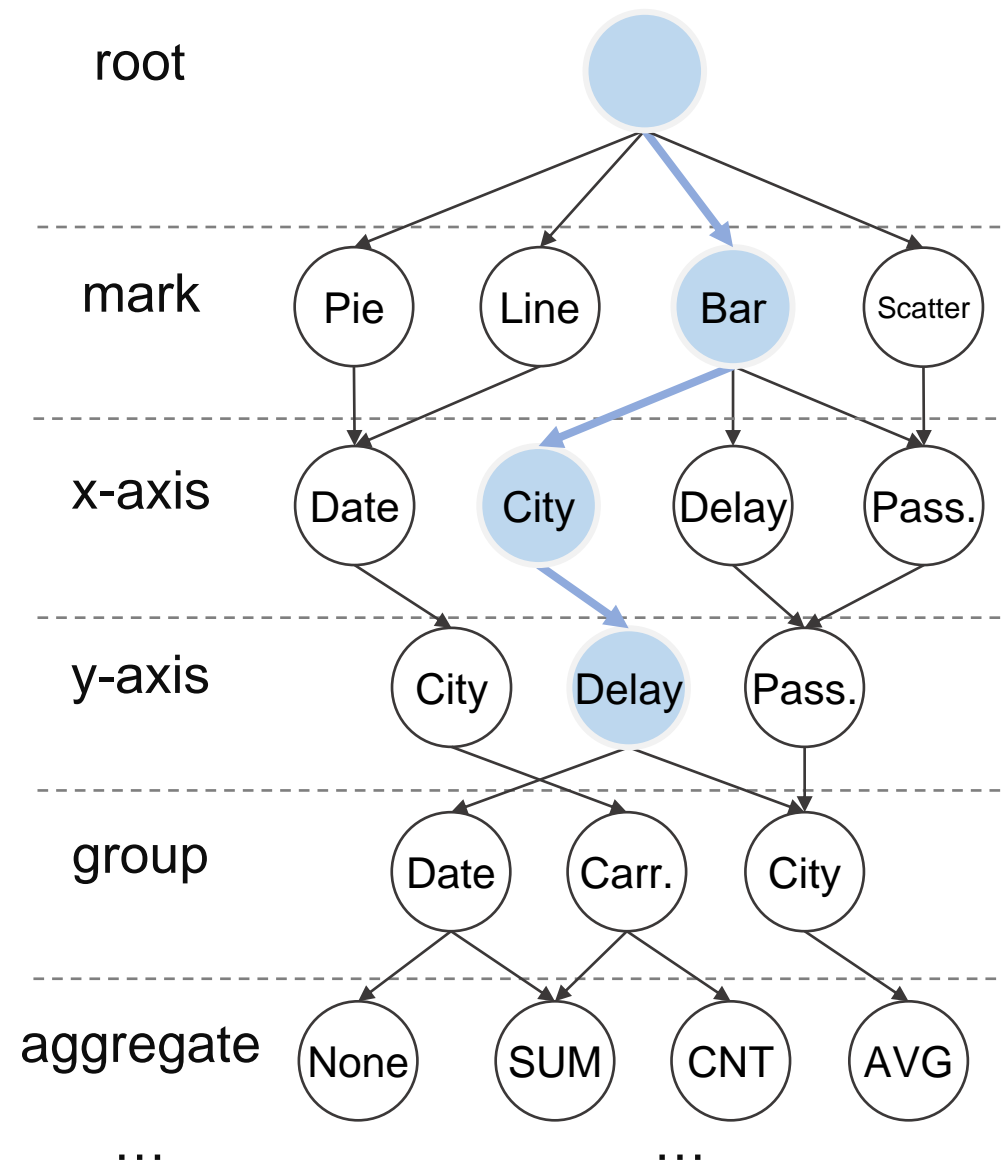
Next, set **City** as the X-axis and **Delay** as the Y-axis.



Current Visualization

Visualization Query Graph

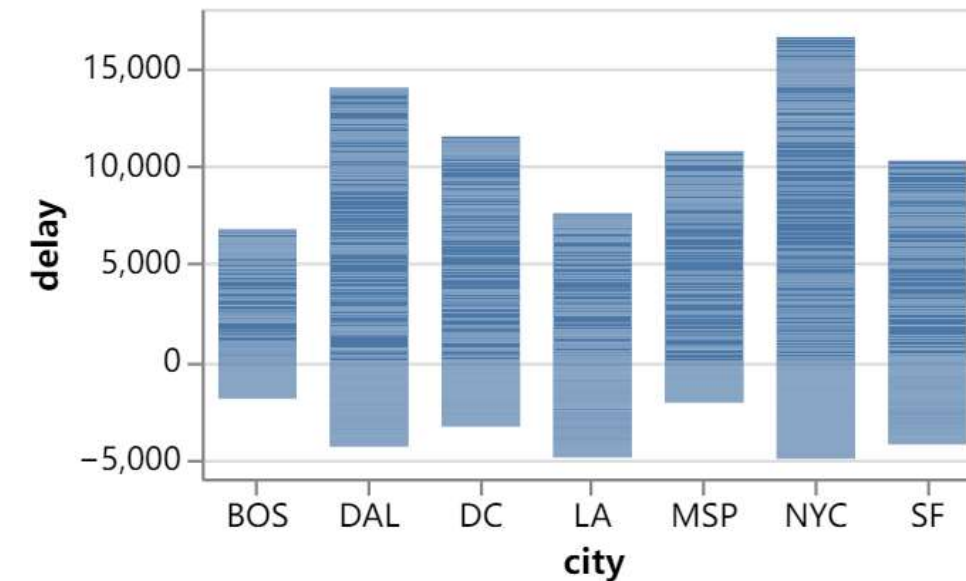
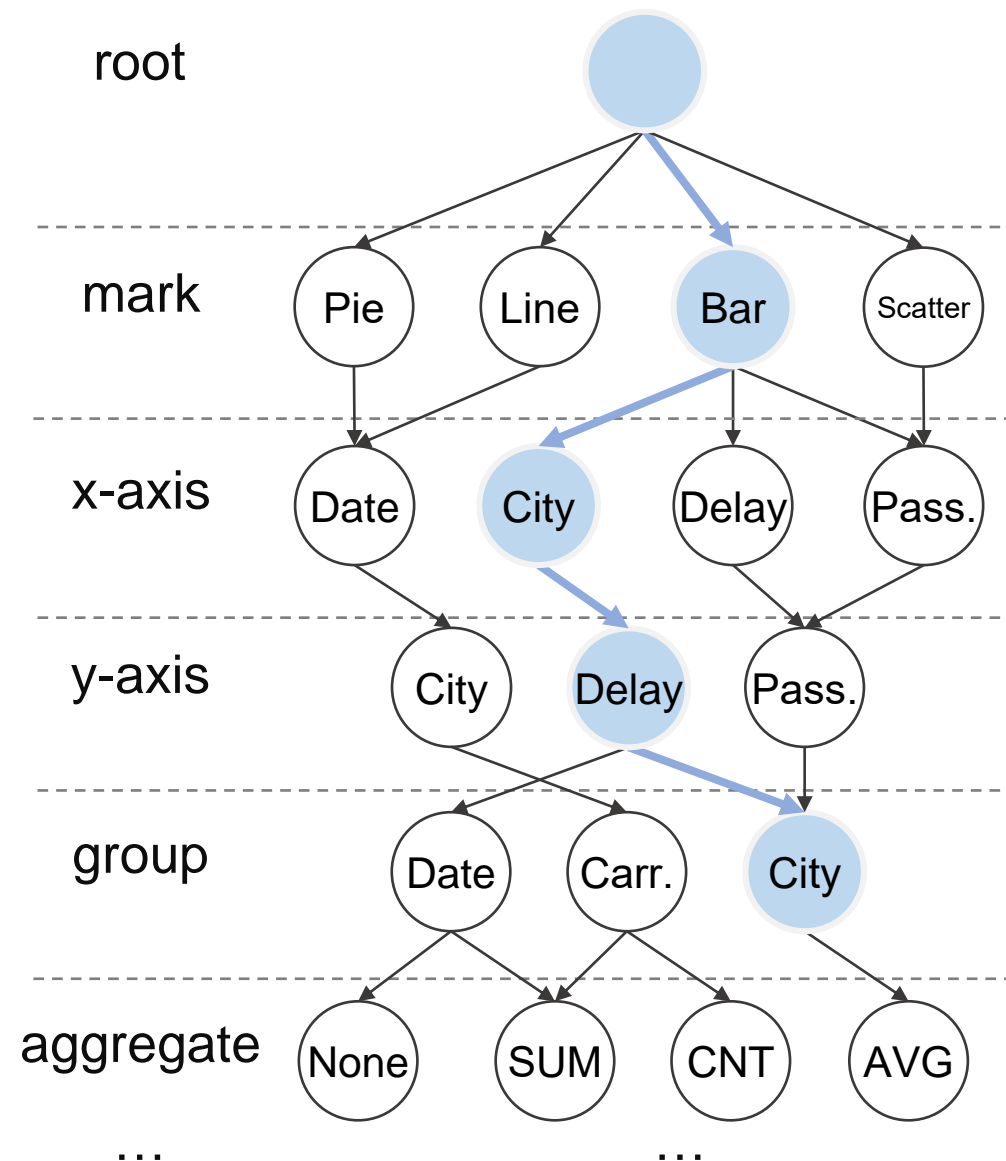
Next, set **City** as the X-axis and **Delay** as the Y-axis.



Current Visualization

Visualization Query Graph


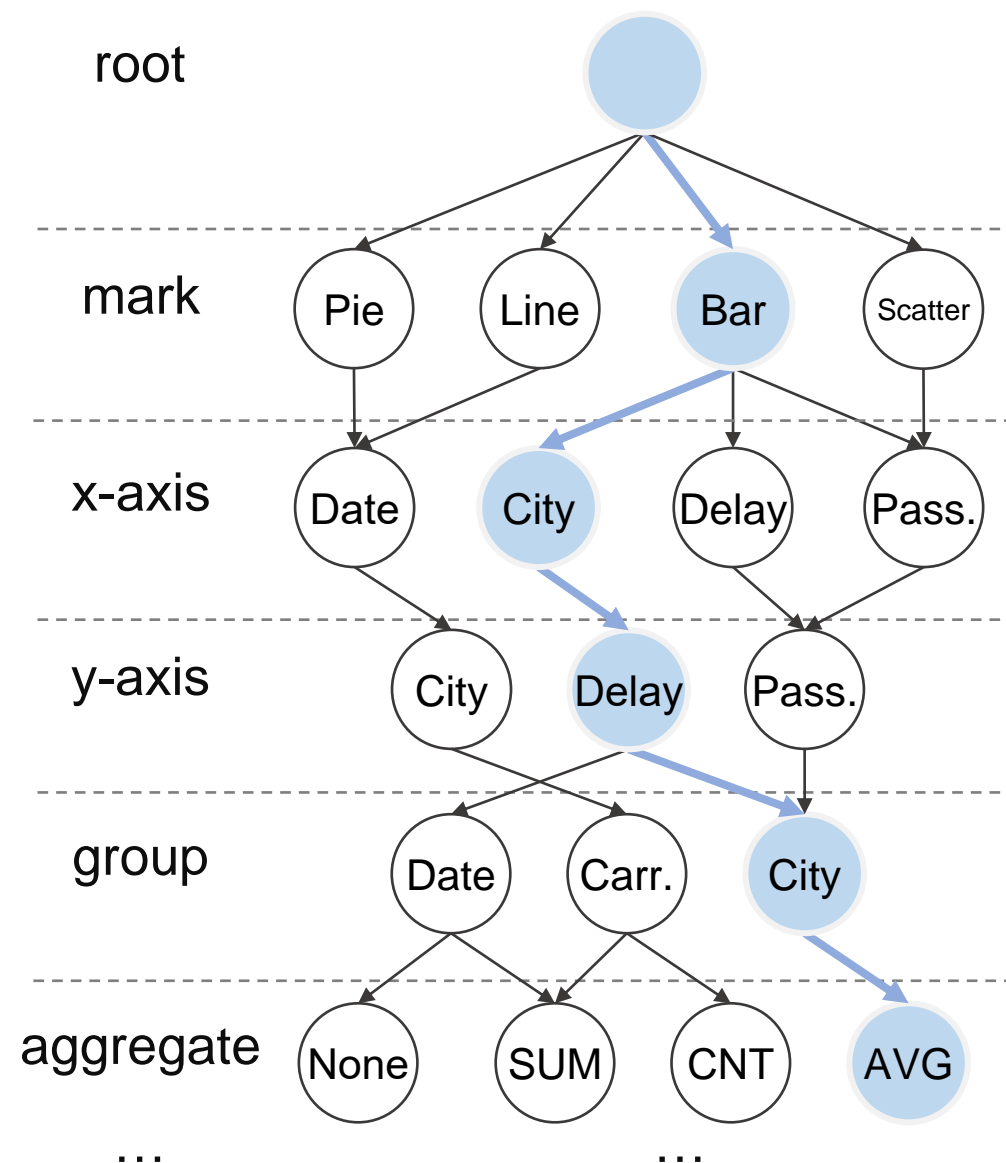
Then, apply a data transformation by **grouping the data by City**.



Current Visualization

Visualization Query Graph

Finally, select the **average** as the aggregation function.

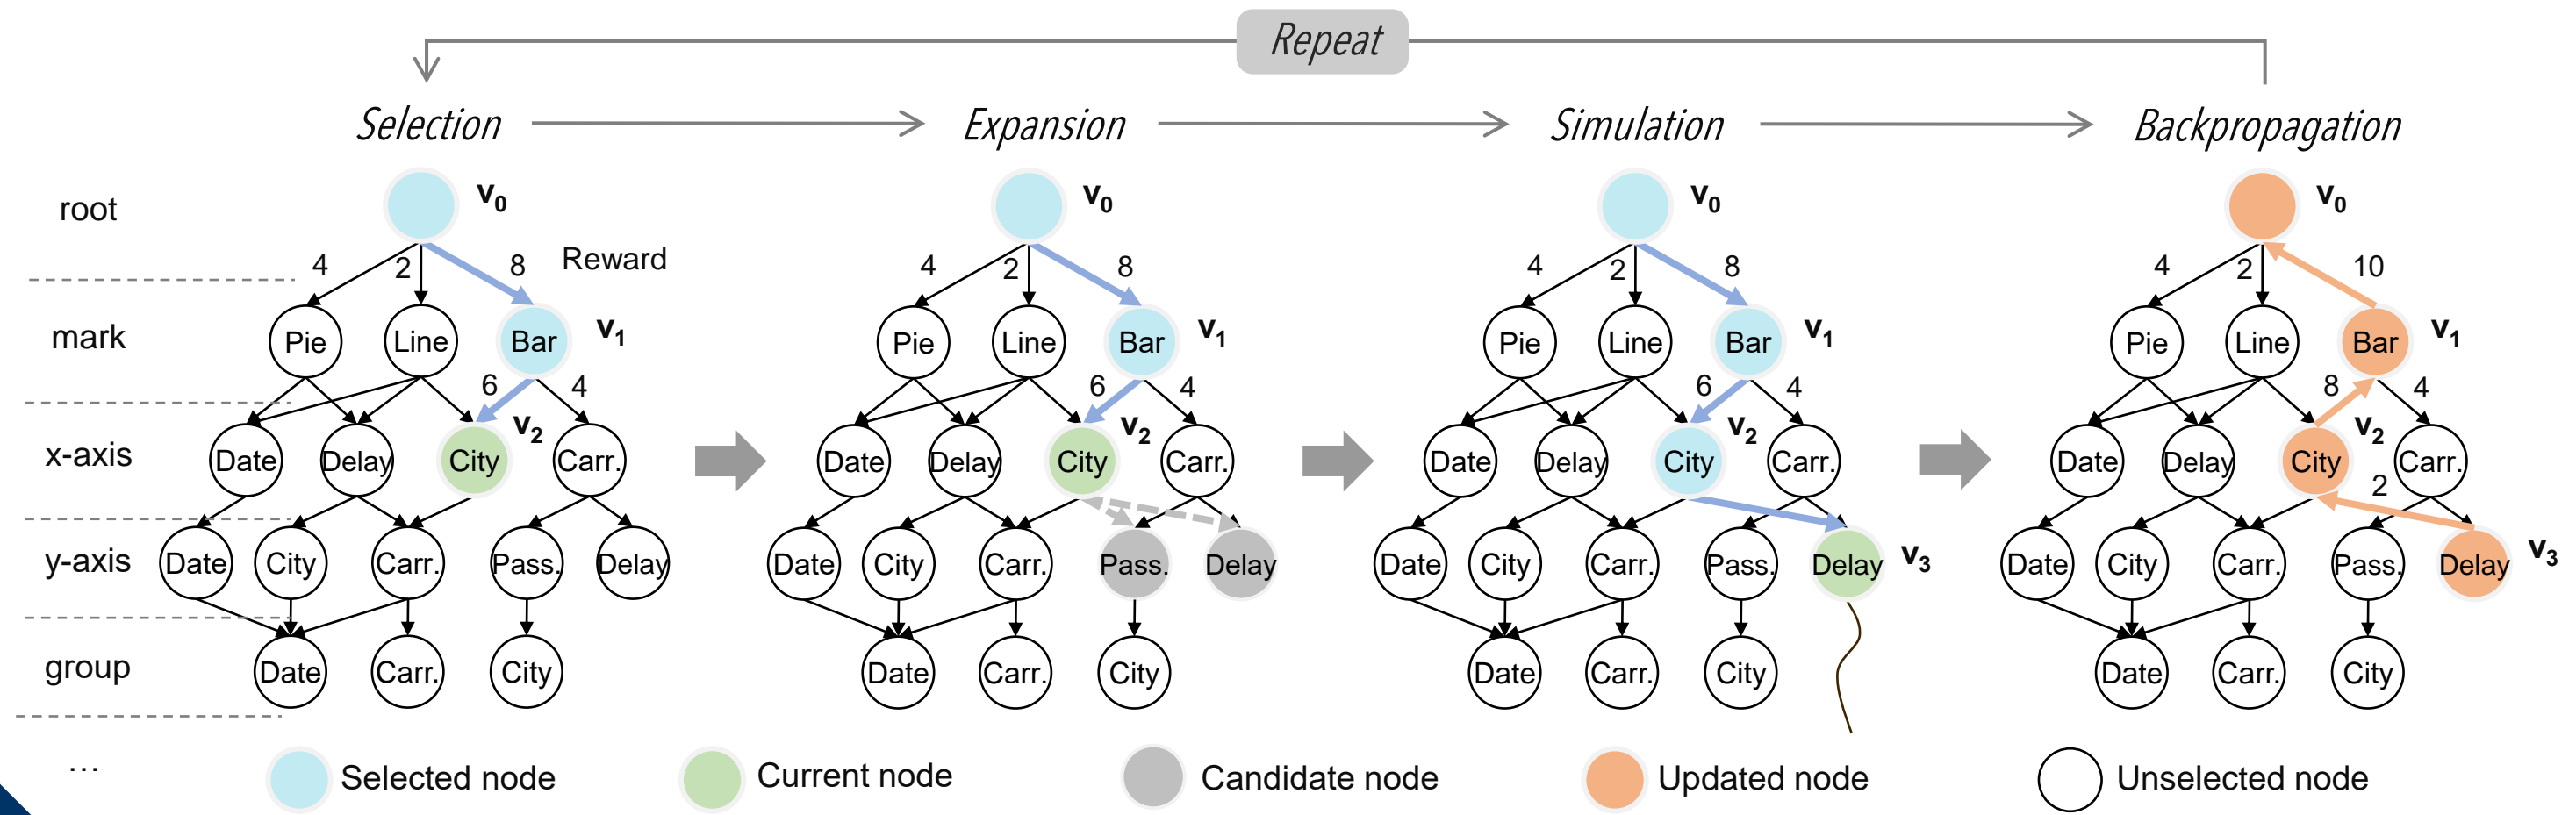


Current Visualization

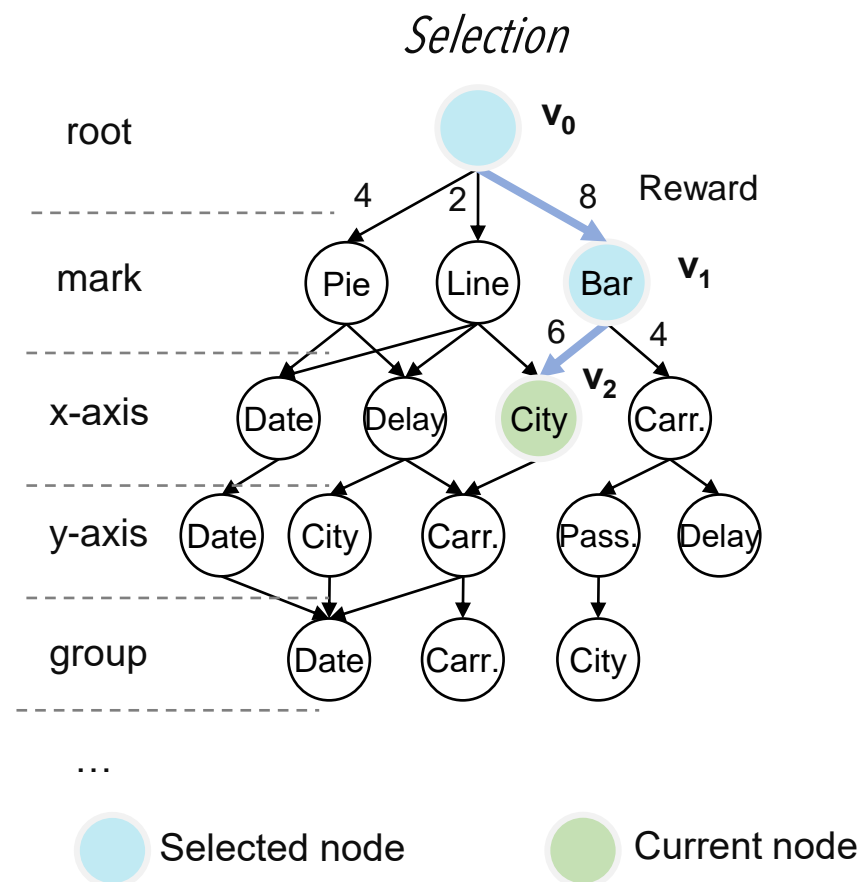
This sequence results in a **valid visualization**.

Monte Carlo Graph Search-based Visualization Generation

The figure shows the MCGS algorithm's **four steps**.



MCGS-based Visualization Generation

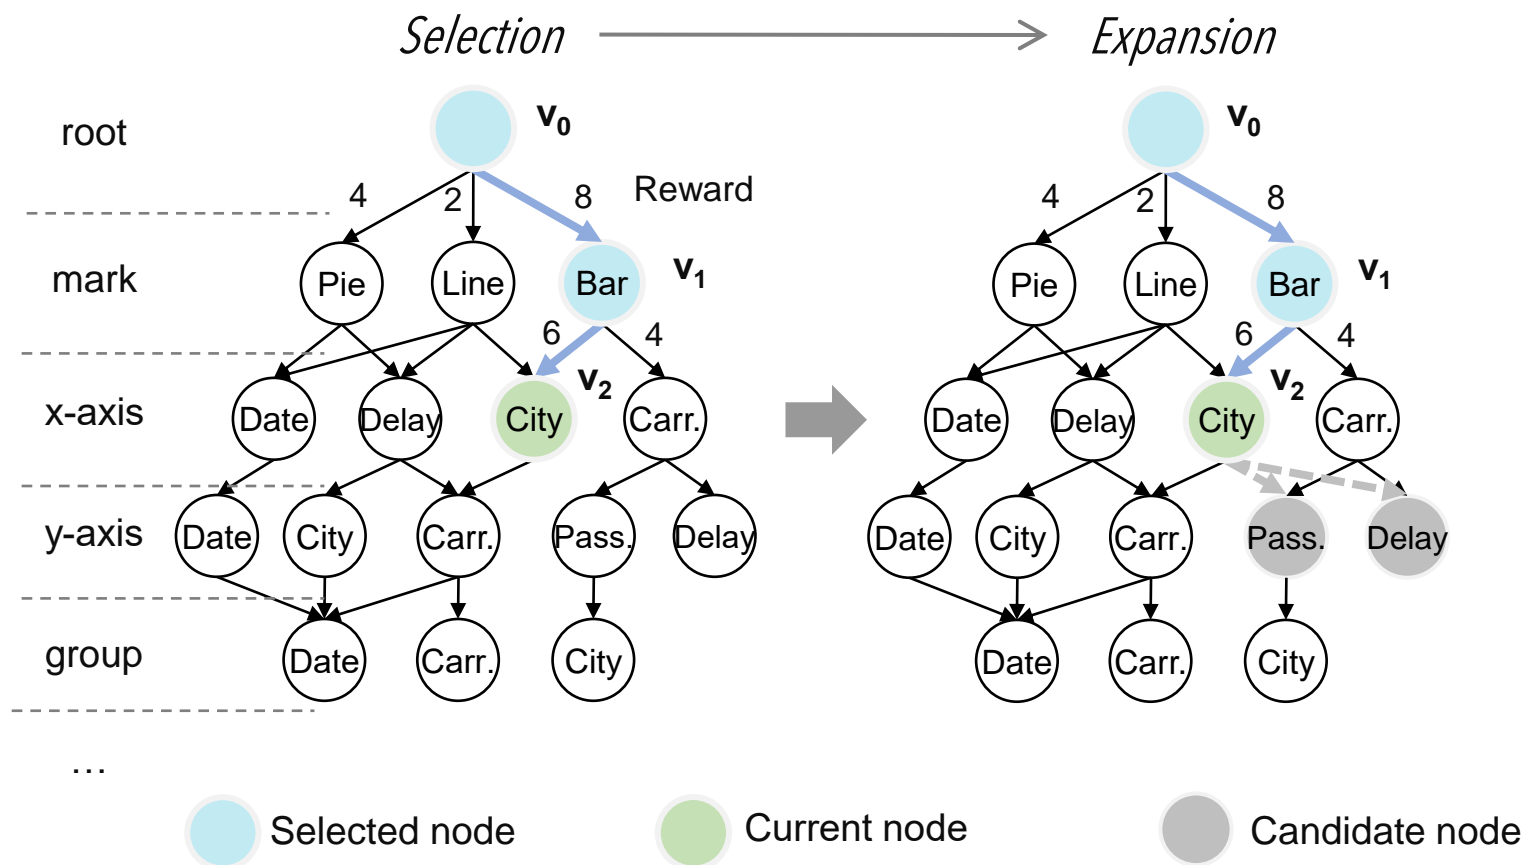


In the **selection** phase, the algorithm uses the **Upper Confidence Bound (UCB)** to recursively **select optimal child nodes** until reaching an unexpanded node.

$$UCB = \underbrace{\bar{X}_i}_{\text{exploitation}} + \underbrace{c\sqrt{2\ln n/n_i}}_{\text{exploration}}$$

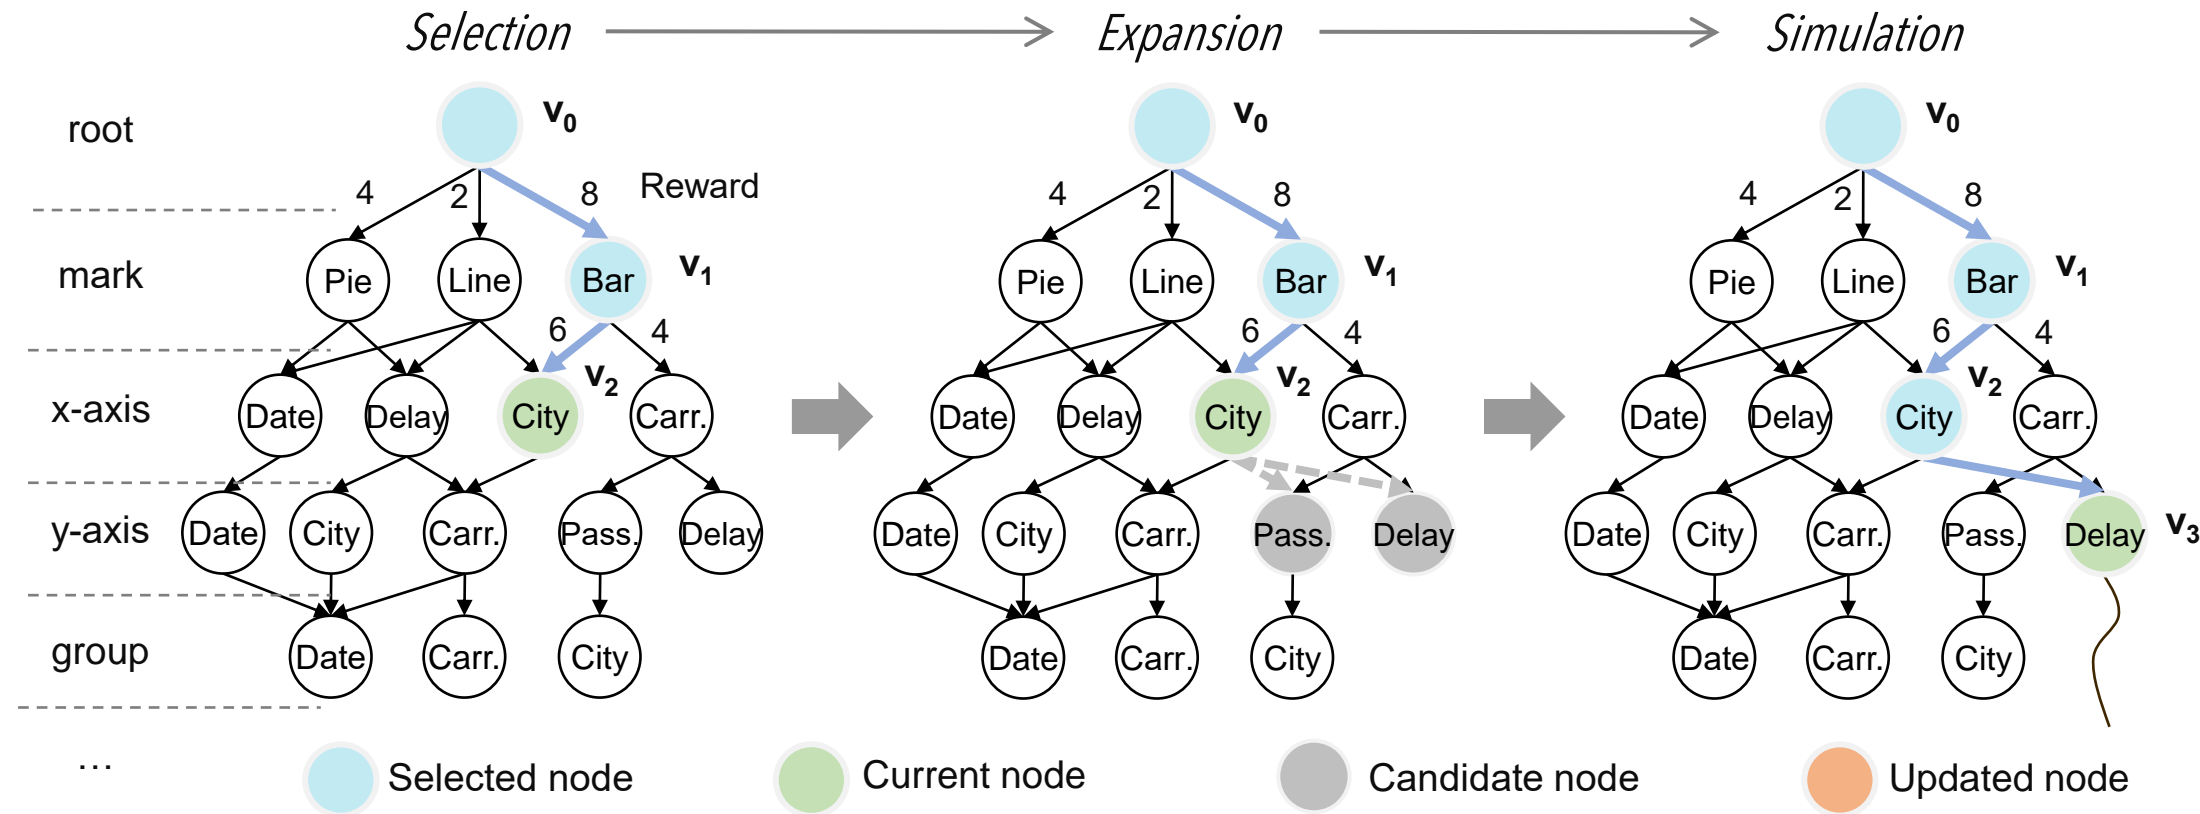
Specifically, the UCB algorithm **balances exploration and exploitation** during the search process.

MCGS-based Visualization Generation



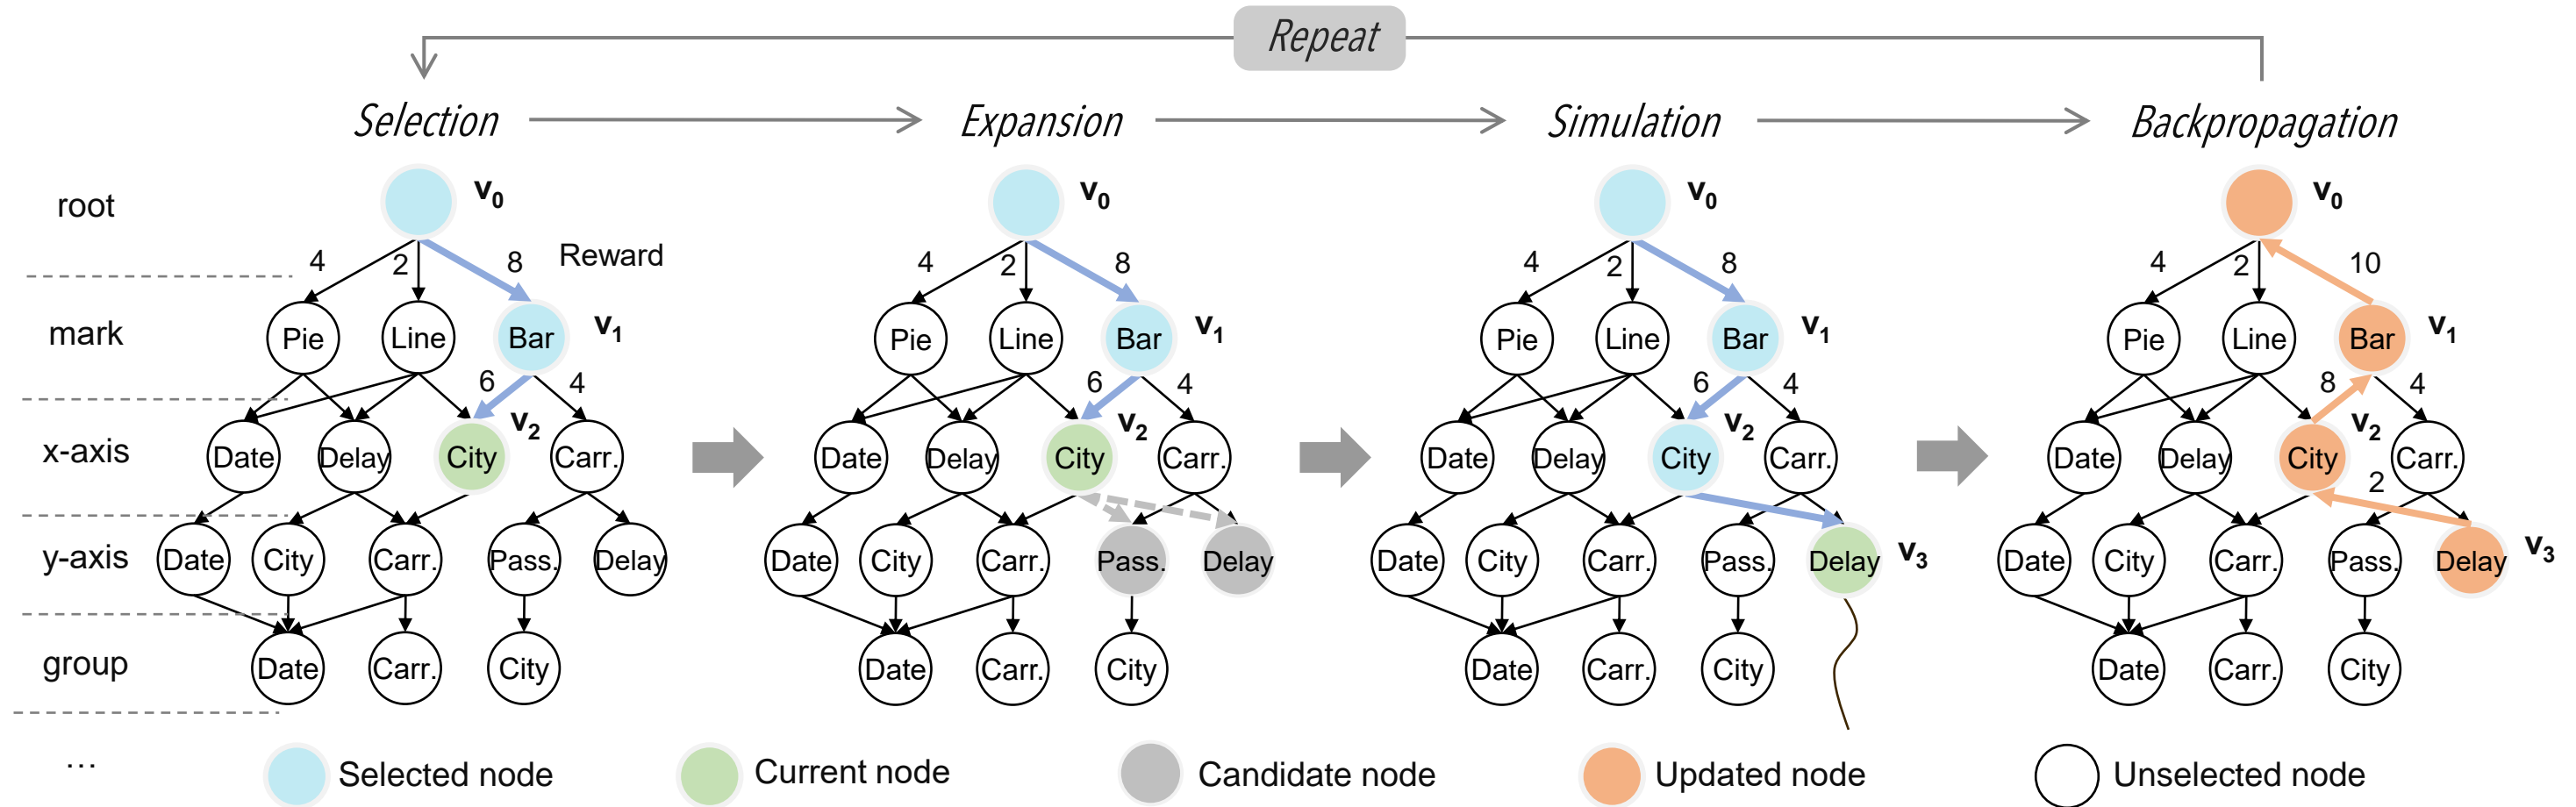
In the **expansion** phase, the algorithm **selects the next valid action** by **removing low-quality visualizations** that are either syntactically incorrect or violate visualization rules.

MCGS-based Visualization Generation



In the **simulation** phase, the algorithm randomly explores based on the current query until a **valid query** is found, then **assigns a reward** using the reward function.

MCGS-based Visualization Generation

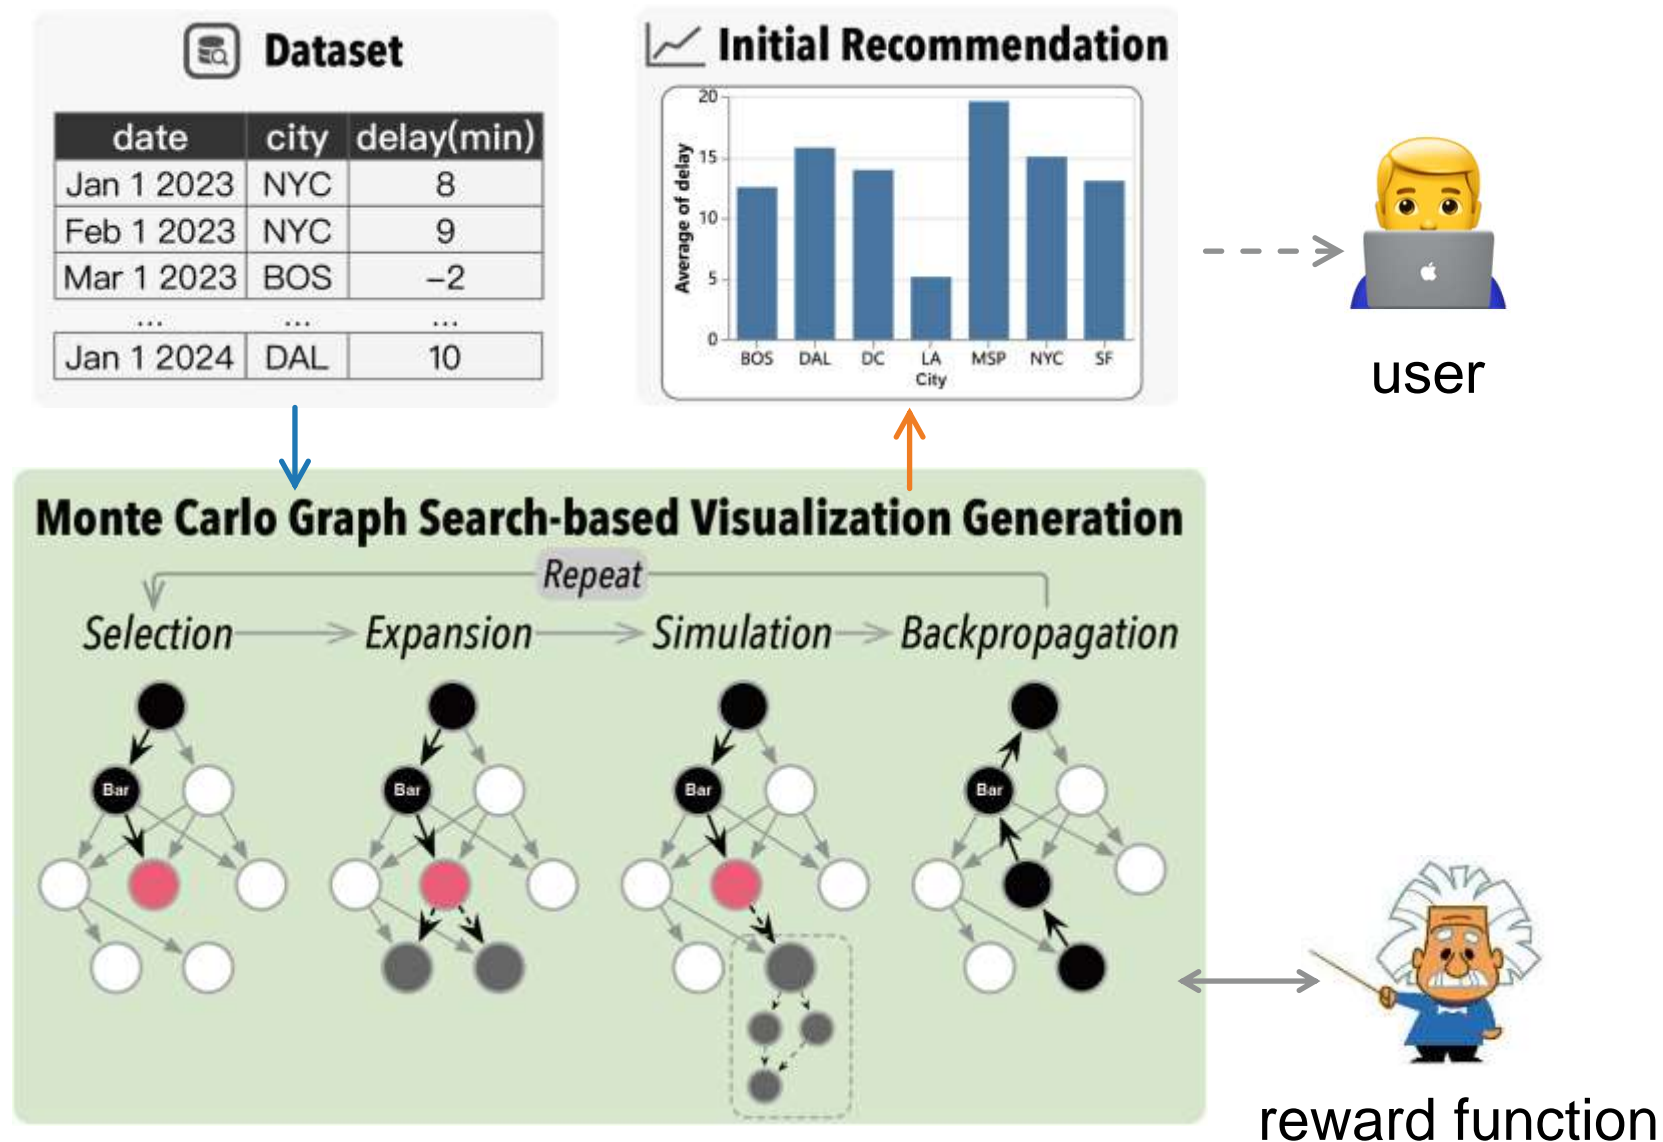


In the **backpropagation** phase, the **reward values** from the simulation are used to **update the graph**, guiding future searches more effectively.

These steps repeat until the maximum iterations are reached.

MCGS-based Visualization Generation


In Monte Carlo Graph Search, the **reward function** plays a key role in guiding the search process.



Challenge 2: How to Evaluate Visualization Quality?

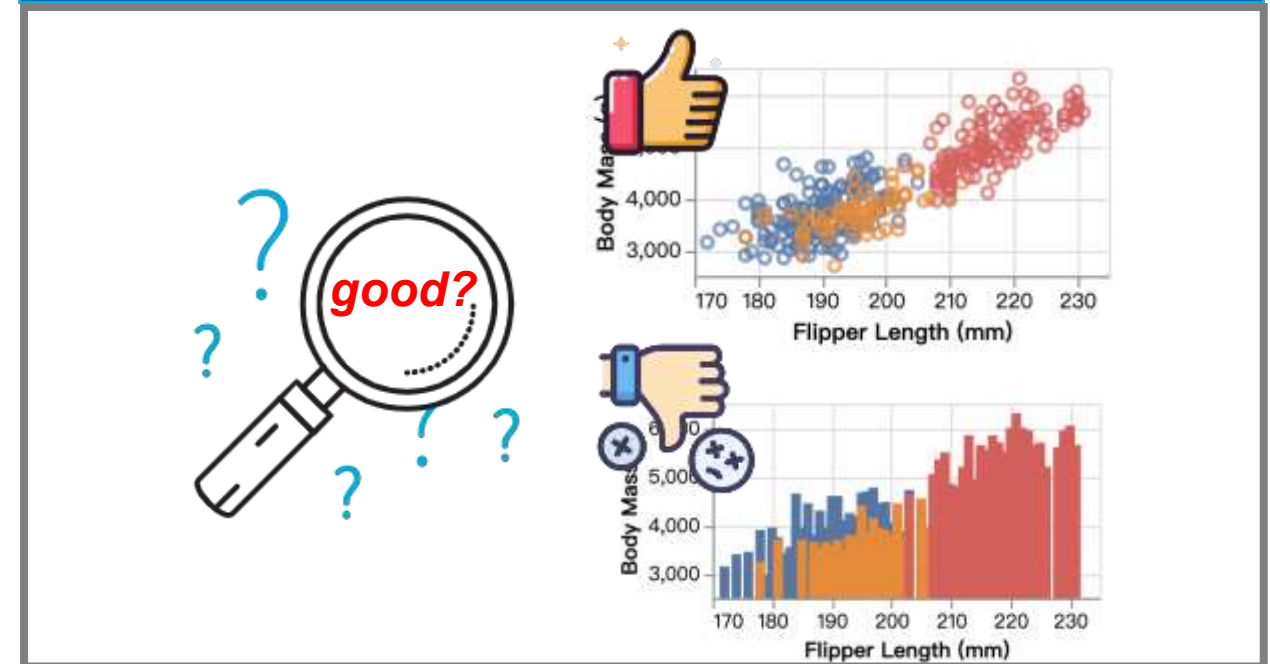
Unlike games such as Go, which have clear rules, **visualization** evaluation **lacks** **well-defined** reward **criteria** and **can be biased** by a **single** metric.

Games have clear rules



Games

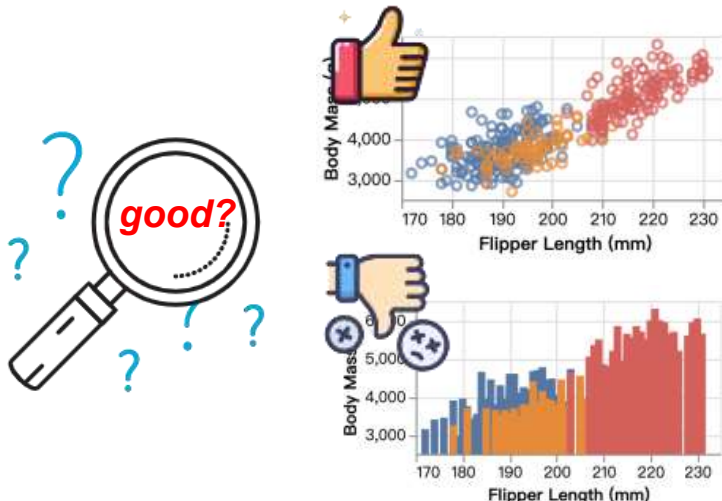
How to quantify the goodness?




Visualization

Composite Reward Function

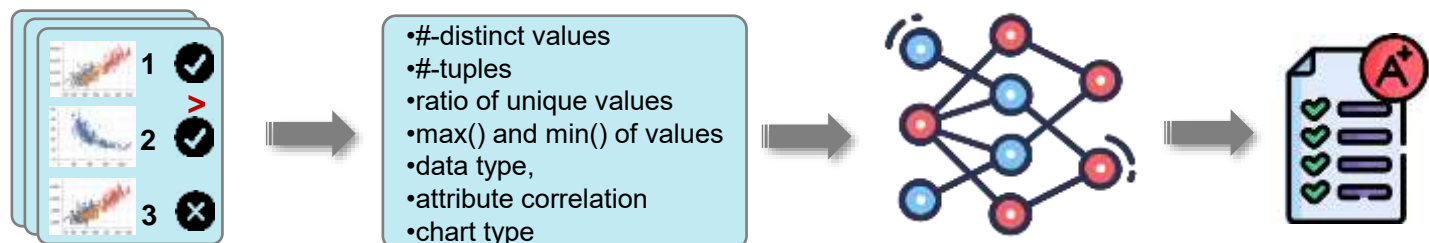
How to quantify the goodness?



Learning from Domain Knowledge



Capturing Data Features



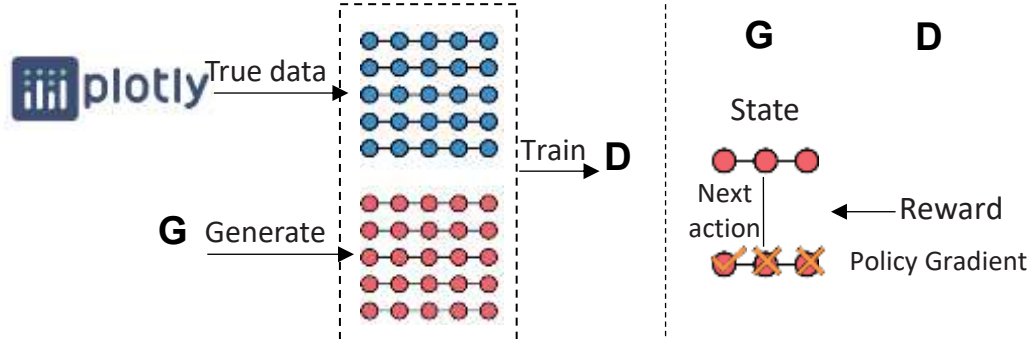
Visualization Corpus

14 types of features

LambdaMART

Scores

Learning Common User Preferences



True data

G Generate

Train

D

State

Next action

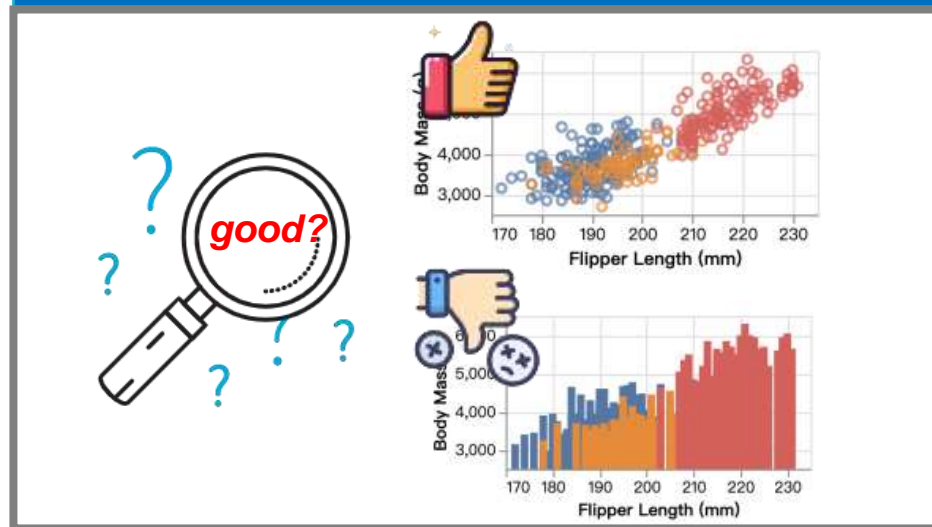
Reward

Policy Gradient


Composite Reward Function

We remove low-quality
visualization results
based on expert rules.


How to quantify the goodness?



Learning from Domain Knowledge



Capturing Data Features



Visualization
Corpus

- #-distinct values
- #-tuples
- ratio of unique values
- max() and min() of values
- data type,
- attribute correlation
- chart type

14 types of features





LambdaMART



Scores

Learning Common User Preferences




Composite Reward Function

How to quantify the goodness?




We extract 14 data features and use LambdaMART to evaluate the visualizations.

Learning from Domain Knowledge



a bar chart with more than 50 bars is hard to get insights.

Capturing Data Features



Visualization Corpus

- #-distinct values
- #-tuples
- ratio of unique values
- max() and min() of values
- data type,
- attribute correlation
- chart type

14 types of features





LambdaMART



Scores

Learning Common User Preferences




Composite Reward Function

How to quantify the goodness?



We use **GANs** to learn common user preferences from the **Plotly** community.


Learning from Domain Knowledge



a bar chart with more than 50 bars is hard to get insights.



Capturing Data Features



Visualization Corpus

- #-distinct values
- #-tuples
- ratio of unique values
- max() and min() of values
- data type,
- attribute correlation
- chart type

14 types of features





LambdaMART



Scores

Learning Common User Preferences



Composite Reward Function

The **composite reward function (CRF)** is calculated as follows:

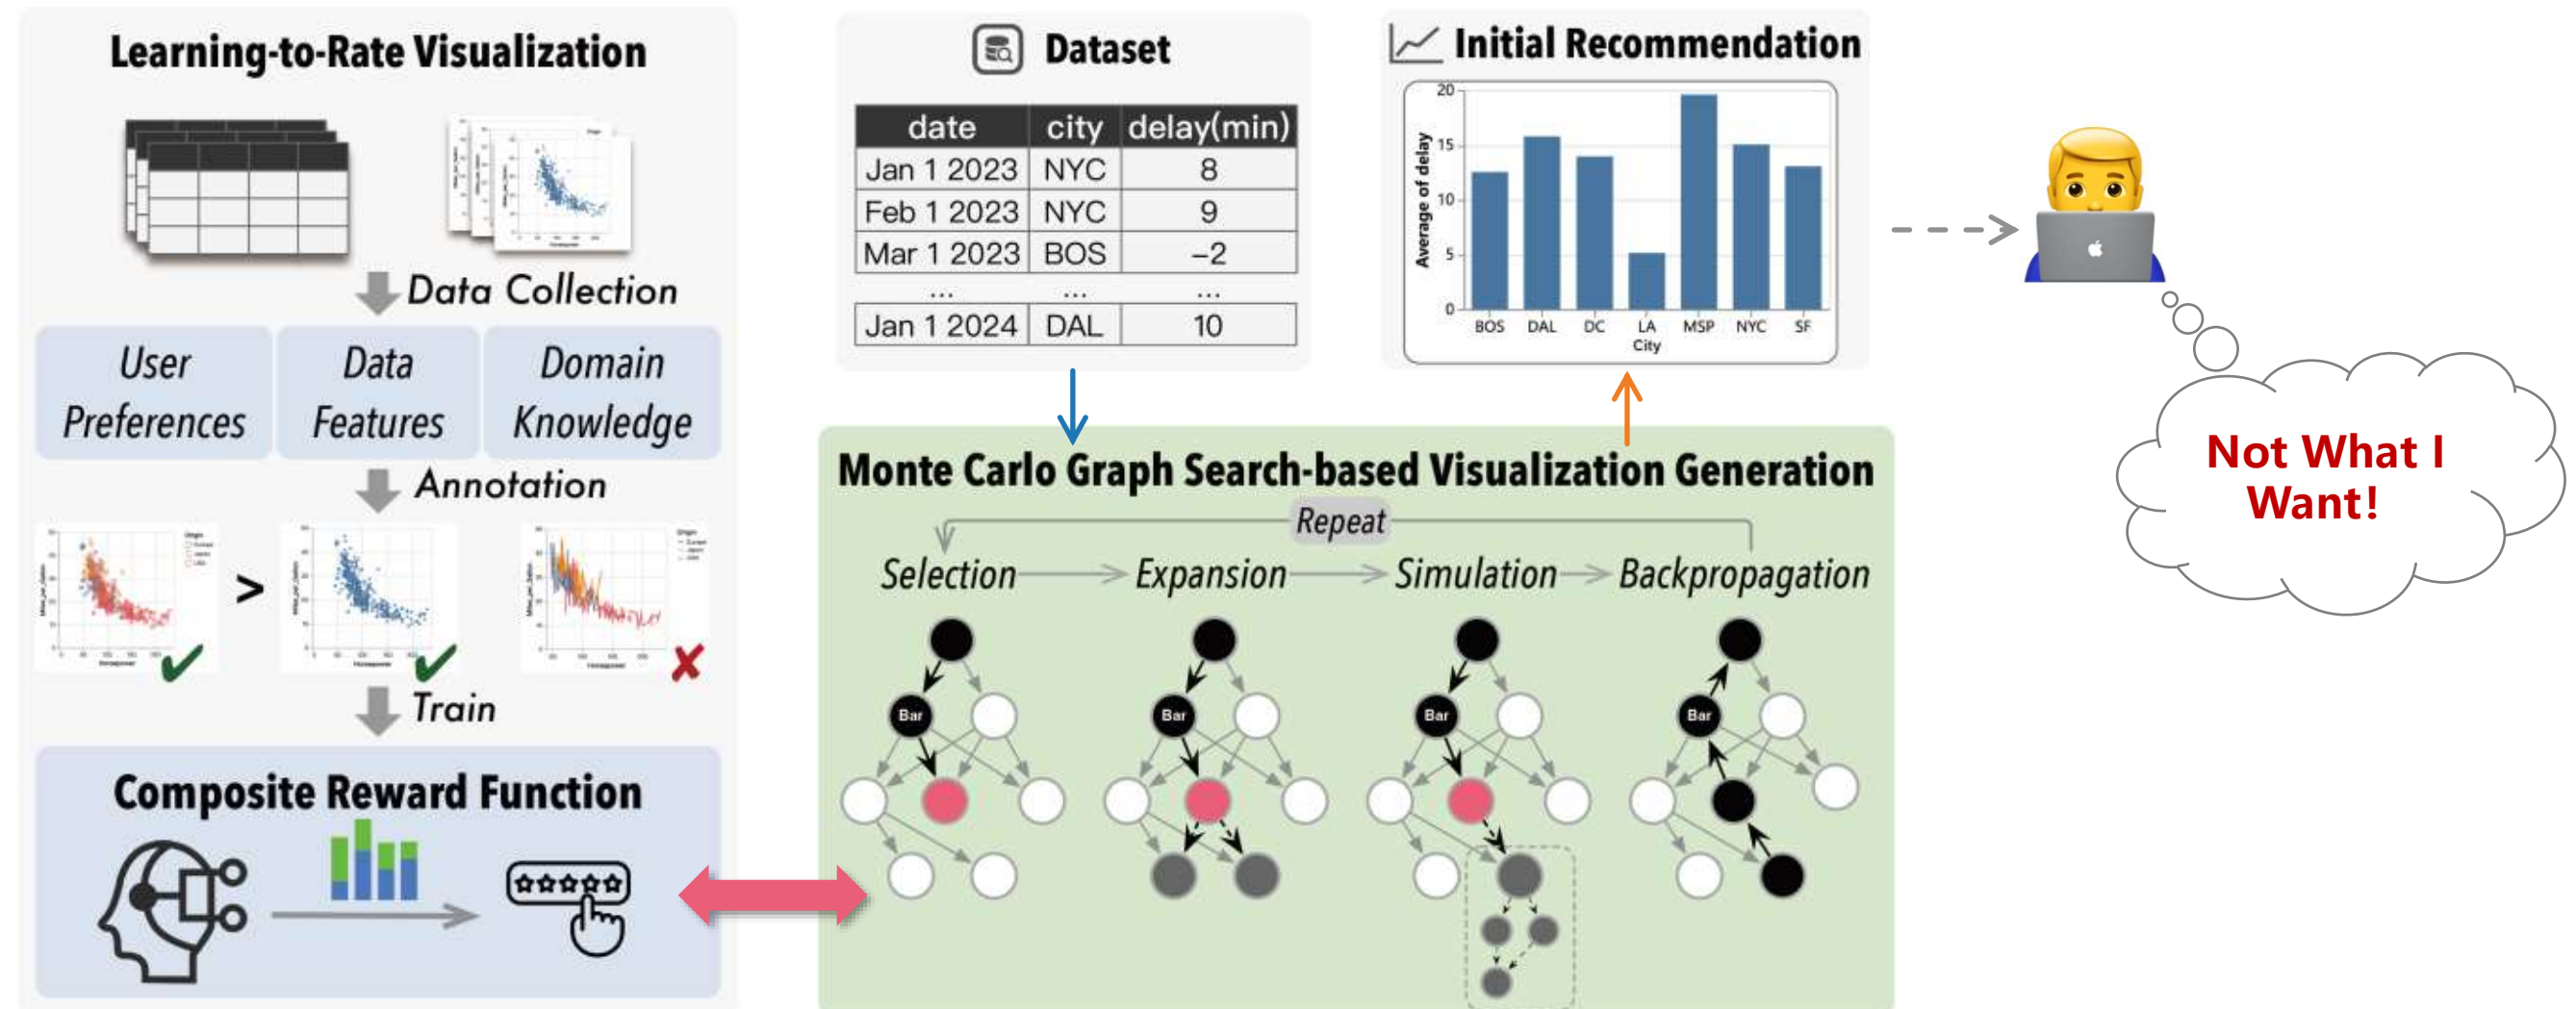
$$CRF = S_k \times (\beta S_d + (1 - \beta) S_u)$$

Where S_k is the domain knowledge score, S_d is the data feature score, and S_u is the user preference score. If S_k is 0, the reward CRF is 0; if S_k is 1, CRF is a weighted combination of S_d and S_u . The weight β controls the importance of S_d and S_u .

Challenge 3: How to Integrate User Feedback?

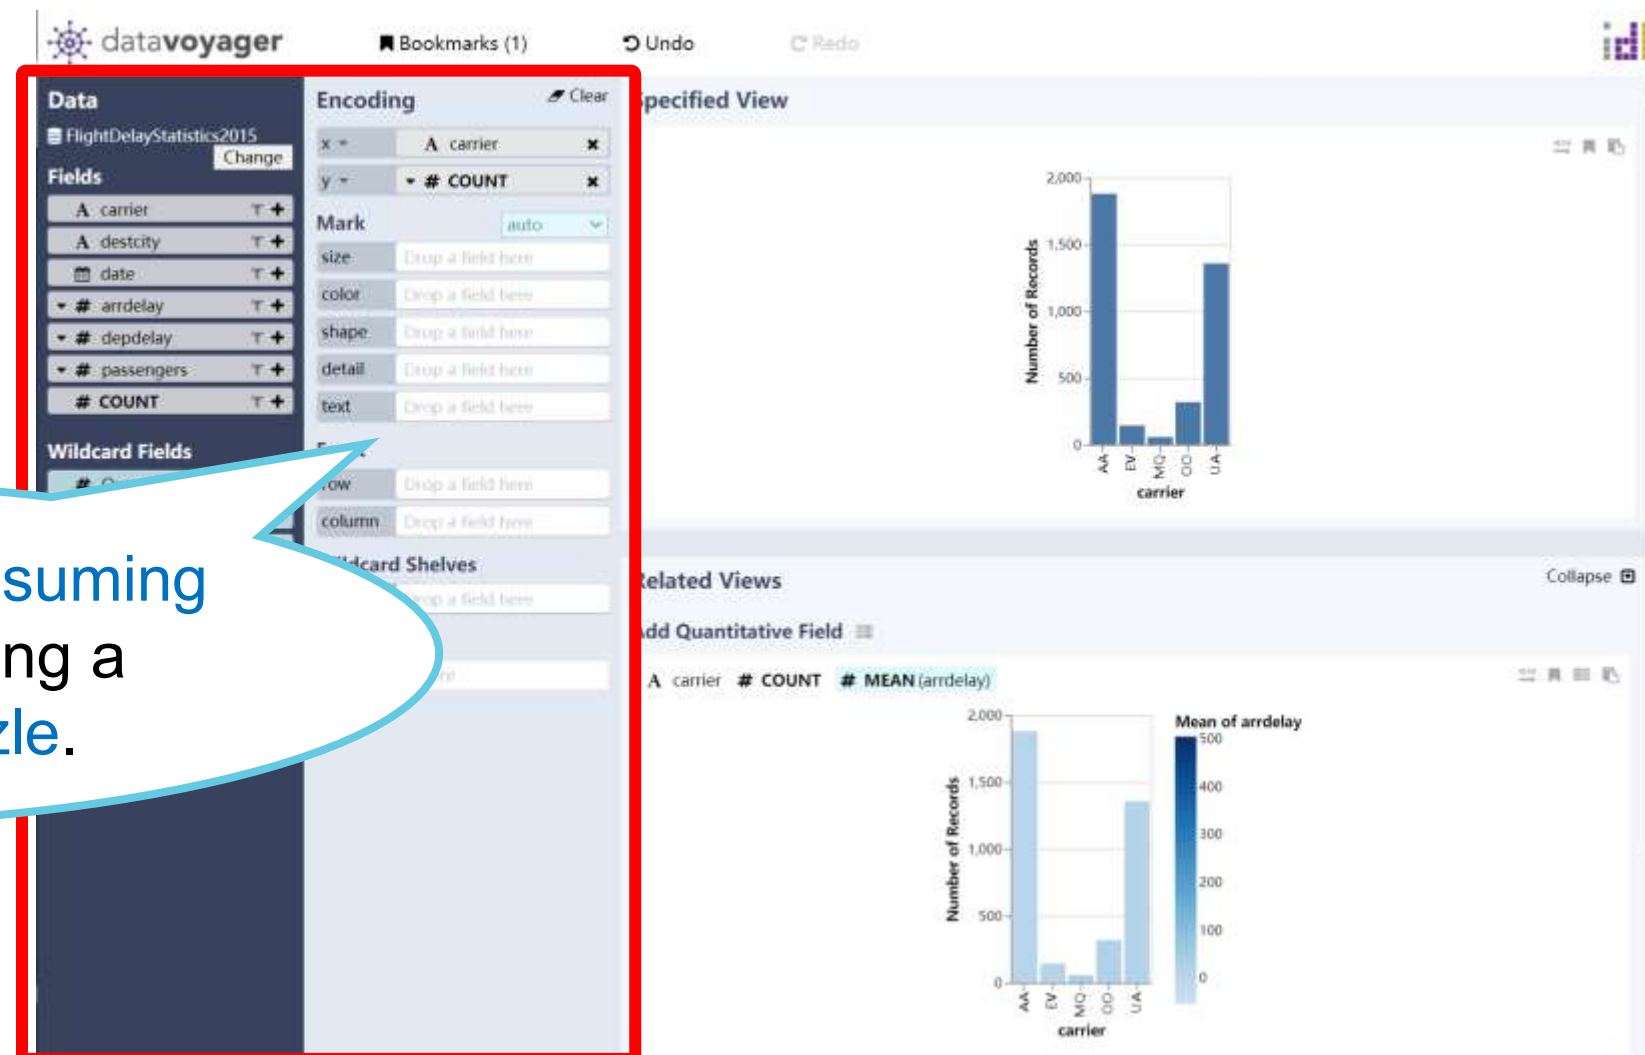
MCGS recommends high-quality visualizations, but they **may not match user needs**.

How can user feedback be used to **align visualizations with user needs**?



How to Integrate User Feedback?

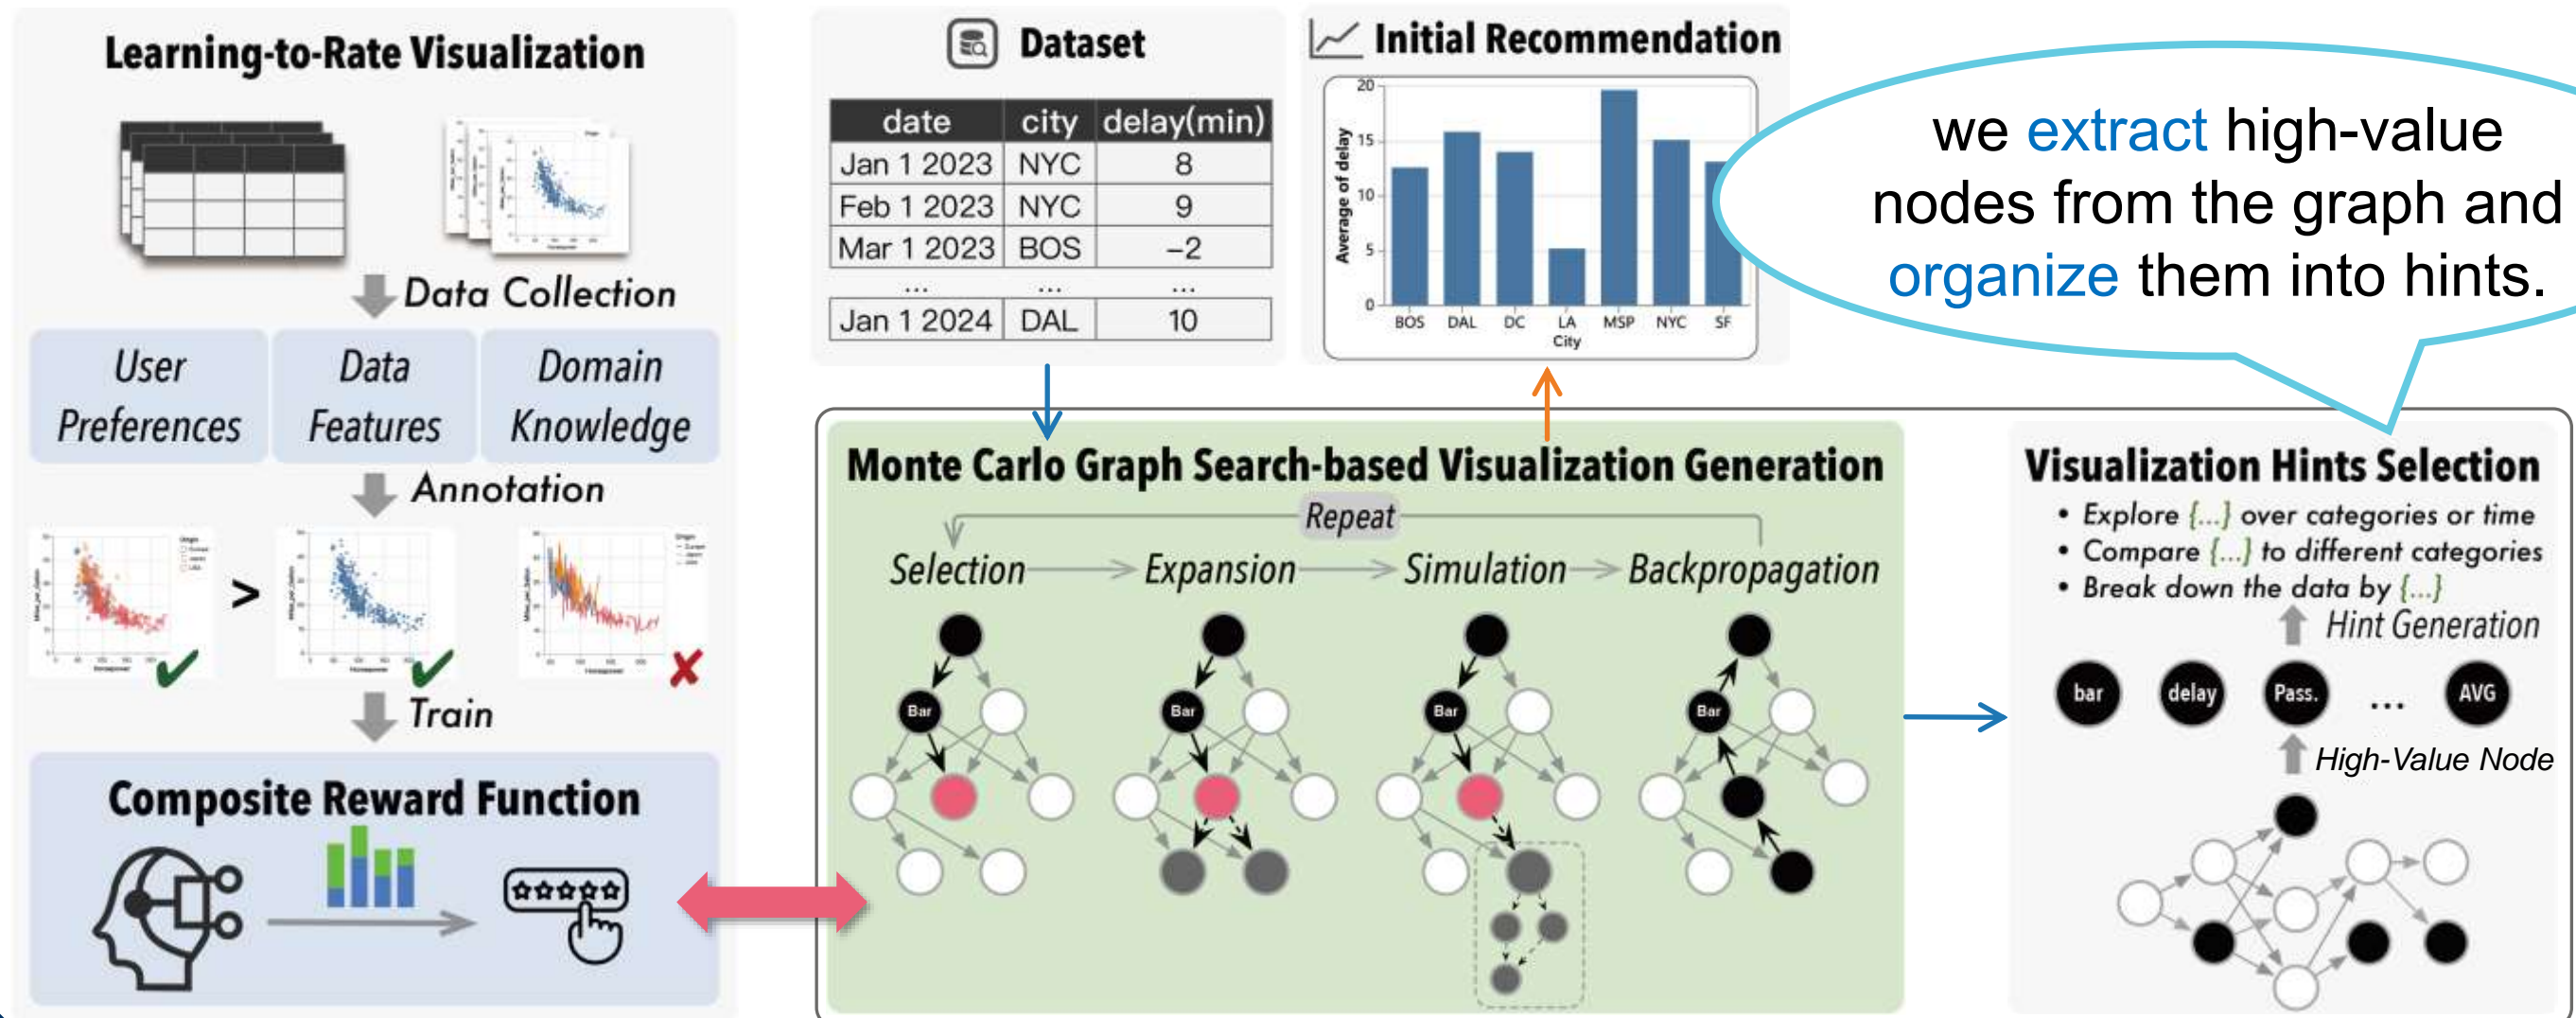
The **common method** is to provide a **control panel** for users to **manually** select data columns and visual encodings.



This process is **time-consuming** and can feel like solving a **fill-in-the-blank puzzle**.

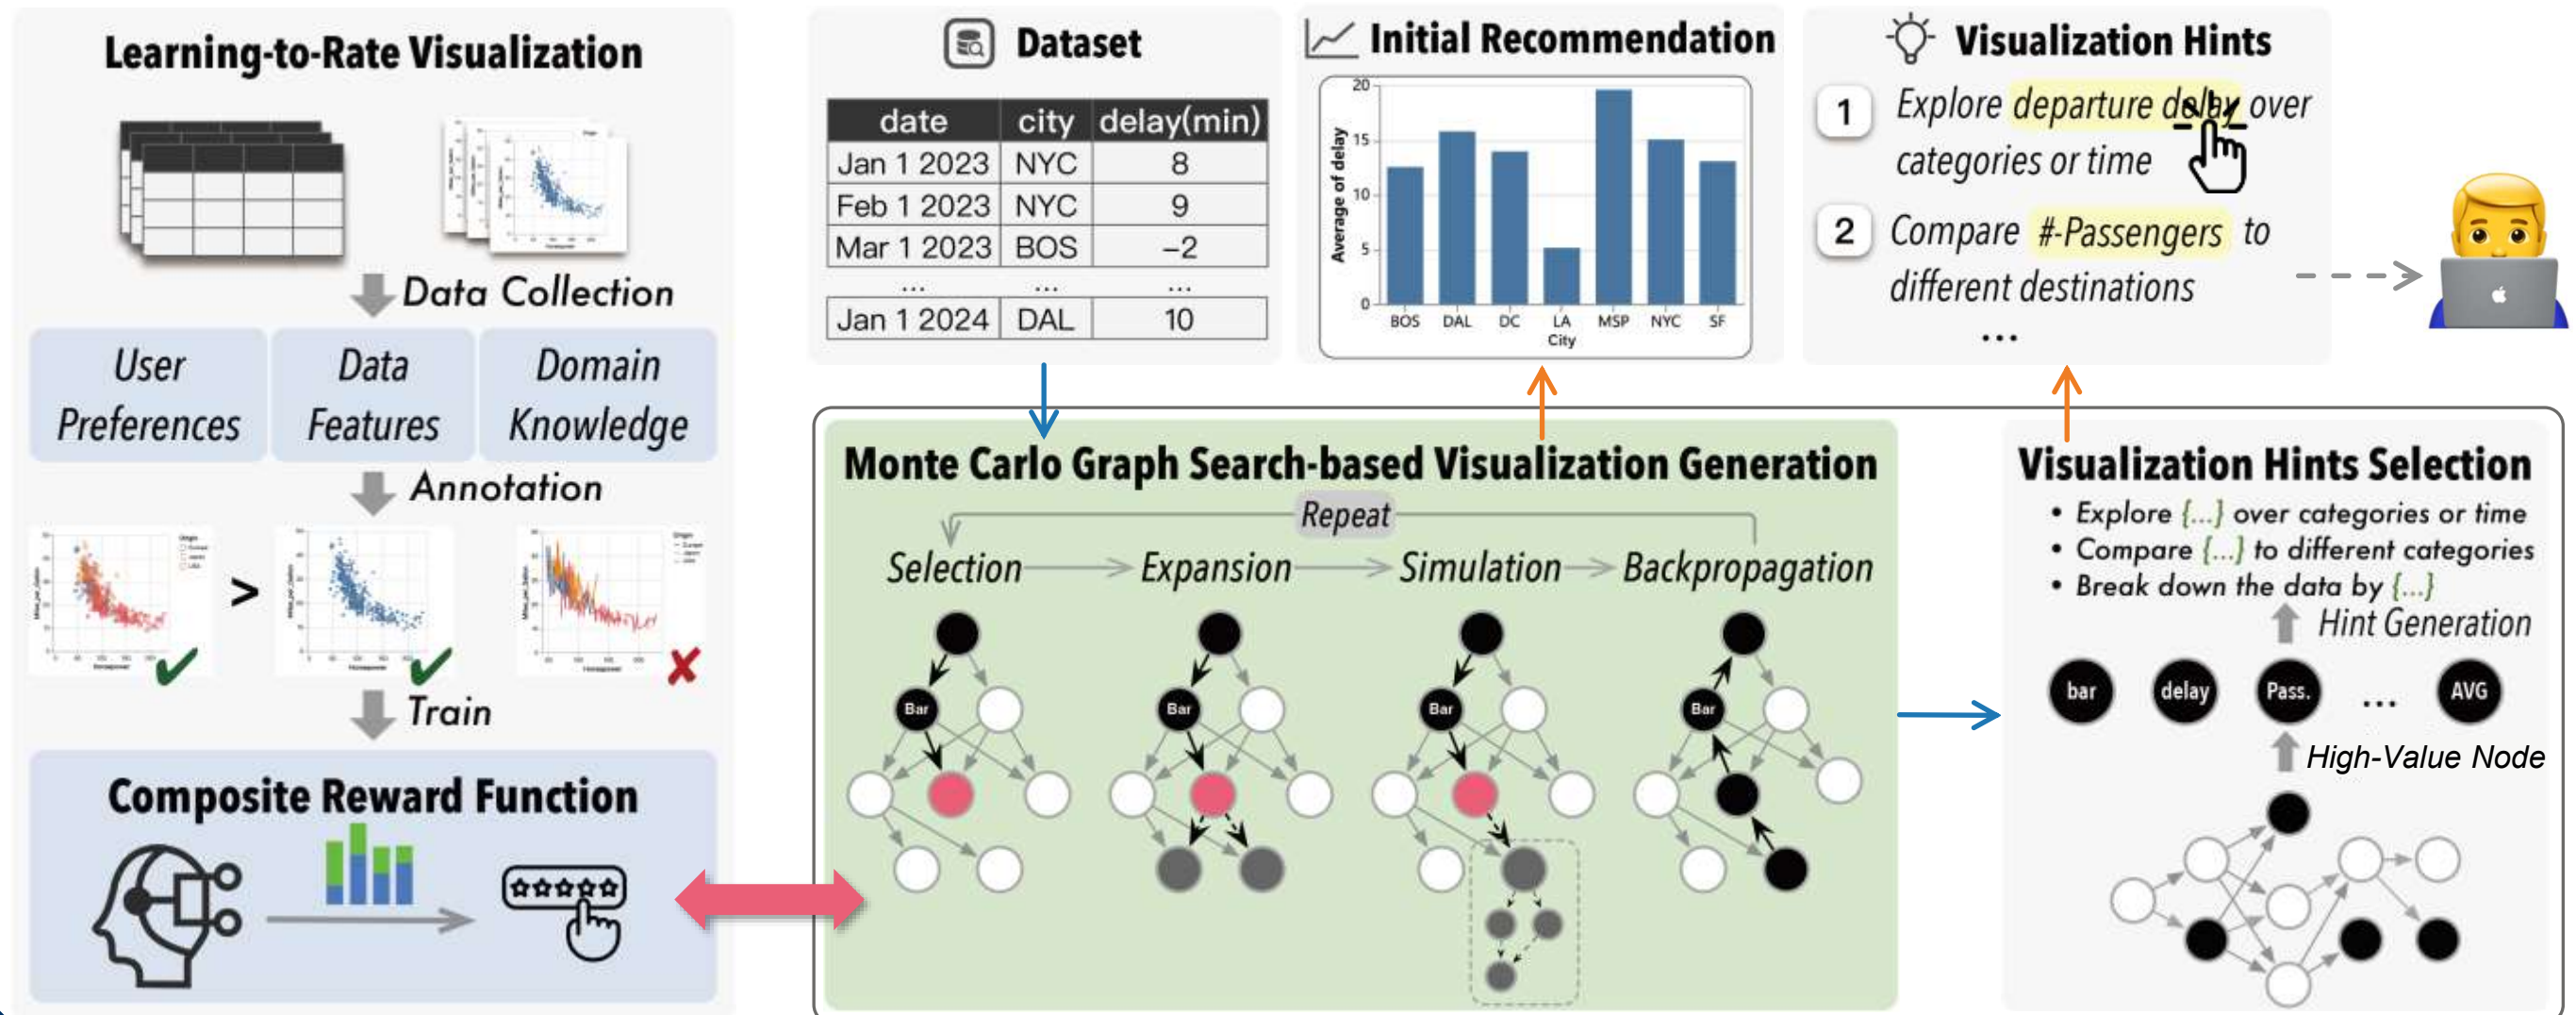
Visualization Hints Module

To address this, we introduce a **visualization hints module**.



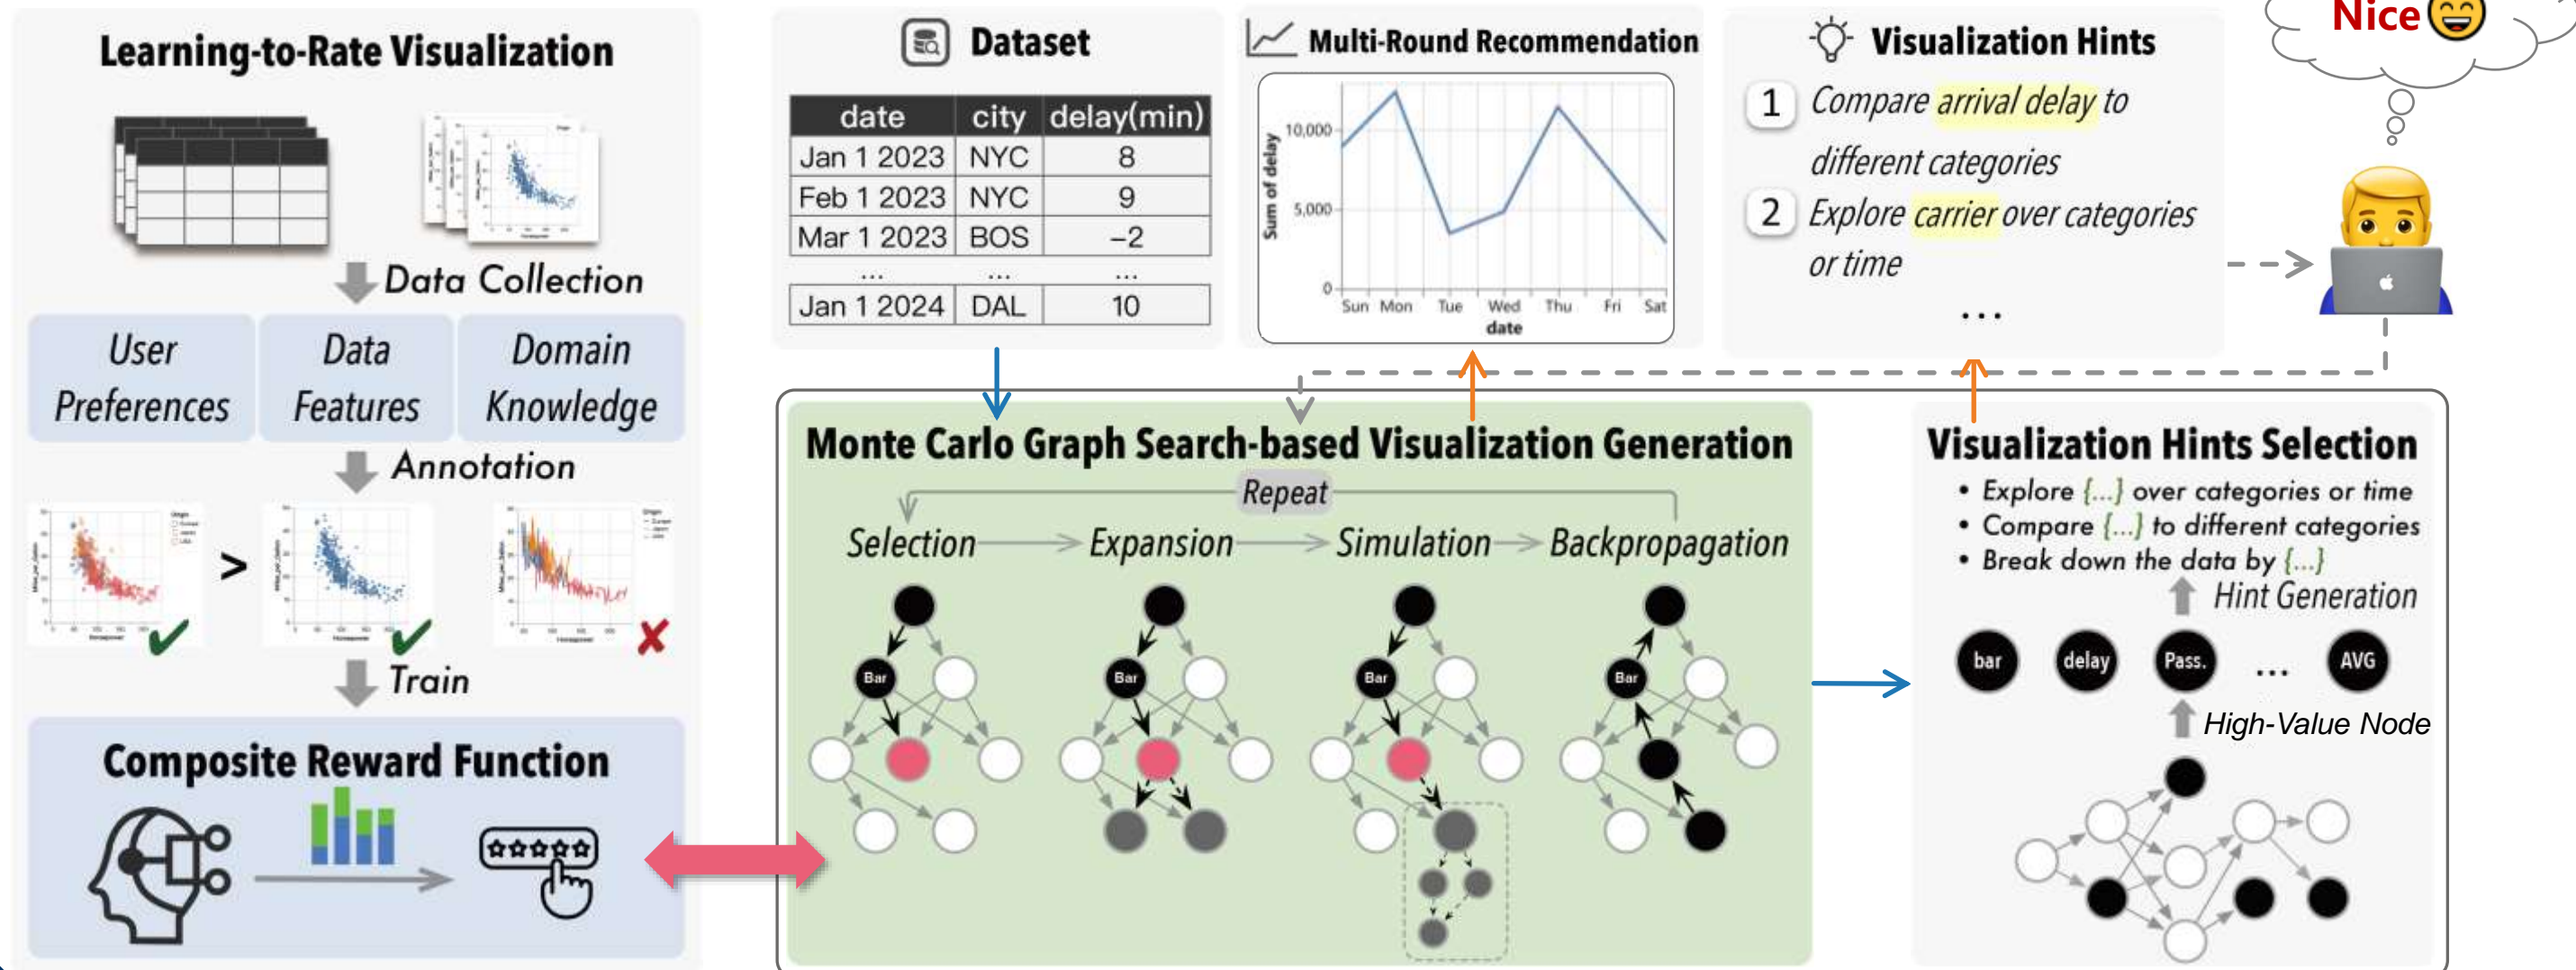
Visualization Hints Module

The user can **indicate** their analysis **intent** by clicking a **hint**.



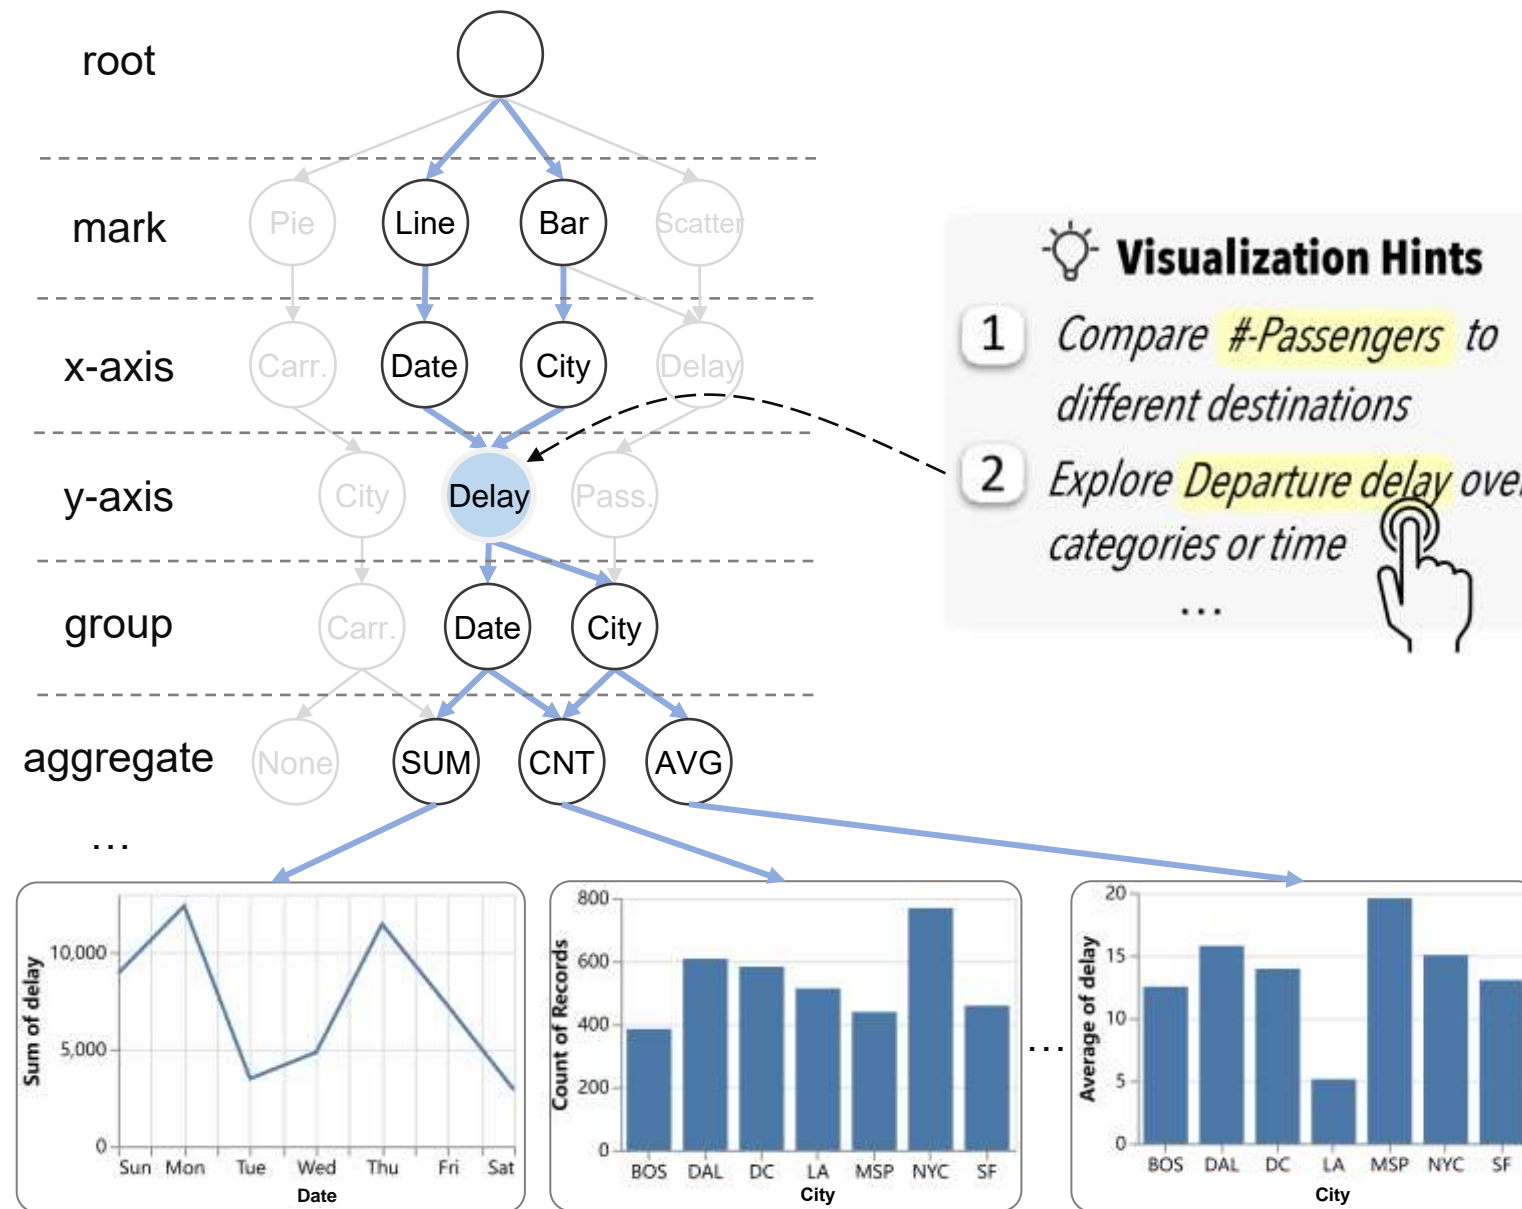
Visualization Hints Module

HAIChart then **updates its recommendations**, offering relevant visualizations and new hints to enhance exploration and **meet user needs**.



Visualization Hint


Visualization Hint: A visualization **hint** corresponds to a **visualization operation** described in **natural language**. Each **hint** is associated with **a set of visualizations**.



Visualization Hint


Candidate hints

Compare **arrdelay** to different categories.



Reward: 3.5
#-Charts: 10

Explore **destcity** across categories or time.



Reward: 2.8
#-Charts: 6

...

Explore **carrier** across categories or time.



Reward: 5.3
#-Charts: 12

Given a set of hints $H = \{h_1, h_2, \dots, h_n\}$, where each hint h_i is associated with a set of visualizations V_i and each visualization $v \in V_i$ has a reward value r_v .

Visualization Hints Selection

We aim to select **top- k hints** that not only **cover different aspects** but also ensure **high-quality** visualizations.

Candidate hints



Reward: 3.5
#-Charts: 10



Reward: 2.8
#-Charts: 6

...



Reward: 5.3
#-Charts: 12

Hints Selection

Top- k hints

- 💡 Explore **carrier** across categories or time.
- 💡 Compare **arrdelay** to different categories.
- 💡 Explore distributions with a **bar** chart.
- 💡 Explore **destcity** across categories or time.
- 💡 Break down the data by **destcity**.
- 💡 Summarize data using **count**.
- 💡 Summarize data using **average**.
- 💡 Compare **depdelay** to different categories.
- 💡 Break down the data by **carrier**.

Visualization Hints Selection

The **problem** is defined as follows:

The goal is to **select** a subset $H' \subseteq H$ of k hints, where the total number of visualizations in H' does **not exceed** a budget B (#-Charts to users), and to **maximize** the overall reward value.

$$\text{Maximize } F(H') = \sum_{h_i \in H'} \sum_{v \in V_i} r_c$$

$$\text{Subject to } \sum_{h_i \in H'} |V_i| \leq B \quad \text{and} \quad |H'| = k$$

This problem is **NP-hard** by a reduction from well-known the **Budgeted Maximum Coverage** problem.

Top- k Hints Selection

To address this, we propose a **top- k visualization hints selection** algorithm.

1. **Select** all hints with costs under budget B to create a **candidate set** (Line 2).
2. **Sort** the hints **based on** their average visualization **score** (Line 3).
3. **Add hints to the final set** until **either** the number of selected hints reaches k **or** the total cost exceeds budget B (Lines 5-10).

Algorithm 2: Top- k Visualization Hints Selection

Input: Set of hints $\mathbb{H} = \{h_1, h_2, \dots, h_n\}$, B , k ;

Output: Selected top- k set of hints \mathbb{H}' ;

```

1  $\mathbb{H}' \leftarrow \emptyset$ ;  $totalCost \leftarrow 0$ ;
2  $\mathbb{H}_v \leftarrow \{h \in \mathbb{H} \mid |h| \leq B\}$ ; // 1. Filter valid hints
3  $\mathbb{H}_v \leftarrow \text{SortByScore}(\mathbb{H}_v)$ ; // 2. Sort hints by score
4 // 3. Selection of top- $k$  hints
5 for each  $h_i$  in  $\mathbb{H}_v$  do
6     if  $|\mathbb{H}'| < k$  and  $totalCost + |h_i| \leq B$  then
7          $\mathbb{H}'.\text{append}(h_i)$ ;
8          $totalCost \leftarrow totalCost + |h_i|$ ;
9     if  $|\mathbb{H}'| = k$  then
10        break;
11 return  $\mathbb{H}'$ ;

```

Experiments Settings

Table 1: Statistics of the experimental datasets (Vis.: Charts)

Datasets	#-Tables	#-Vis.	Avg(#-Vis.)	Avg(#-Rows)	Avg(#-Col.)	Max(#-Col.)
VizML	79,475	162,905	2	2,817.8	3.3	25
KaggleBench	8	252	31.5	32,585.9	9.1	15

Datasets: We used [VizML](#) and [KaggleBench](#) benchmarks.

Evaluation Metrics: We used [Hit@k](#), [P@k](#), and [Rt@k](#).

Comparison Methods: We compared HAIChart with [10](#) state-of-the-art methods: DeepEye, Data2Vis, VizGRank, PVisRec, VizML, LLM4Vis, Voyager2, HAIChart-, LLM4Vis+, and MCTS-based baseline.

Effectiveness of the First-round of Recommendations

D	Tasks	Metrics	The State-of-the-Art Methods						Our Methods		
			Data2Vis [10]	VizGRank [11]	DeepEye [20]	PVisRec [24]	VizML [14]	LLM4Vis [39]	MCTS	HAICHART-	HAICHART
VizML	Data Queries	Hit@1	47.5%	57.6%	52.4%	52.3%	-	-	78.3%	79.7%	79.3%
		Hit@3	51.3%	67.2%	67.6%	58.7%	-	-	88.2%	91.3%	91.9%
	Design Choices	Hit@1	41.7%	34.9%	34.1%	28.9%	28.7%	47.9%	42.4%	50.6%	48.7%
		Hit@3	43.7%	42.9%	40.7%	51.3%	-	-	77.1%	81.8%	81.5%
	Overall	Hit@1	24.3%	25.6%	25.7%	21.8%	-	-	33.1%	37.9%	36.9%
		Hit@3	26.9%	30.1%	33.9%	42.3%	-	-	64.7%	68.4%	67.4%
KaggleBench	Data Queries	P@10	41.2%	58.7%	62.5%	42.5%	-	-	52.2%	60.0%	63.8%
		R10@30	25.0%	50.0%	48.7%	67.5%	-	-	73.6%	80.1%	83.7%
	Design Choices	P@10	88.7%	87.5%	93.7%	91.9%	Hit@2:78.3%	Hit@2:87.6%	93.8%	96.3%	96.3%
		R10@30	95.0%	81.3%	95.0%	85.0%	-	-	92.5%	96.2%	96.2%
	Overall	P@10	28.7%	43.7%	48.7%	36.7%	-	-	45.4%	51.3%	55.0%
		R10@30	13.8%	41.3%	33.7%	60.0%	-	-	63.8%	72.5%	74.9%

Overall, our methods (HAICHART and HAICHART-) significantly **outperform all state-of-the-art methods across all metrics**, showing the effectiveness of our framework.

Effectiveness of Multi-round Recommendations

We design a **user study**:

- **Participants:** 17 participants. 12 experts and 5 non-experts.
- **Task:**
 - Participants used HAIChart, DeepEye, LLM4Vis+, and Voyager2 with **KaggleBench** datasets.
 - Tasks included **specific analyses** and **open-ended** explorations.
- **Procedure:**
 - Preparation: **Introduced** study context, datasets, and tools.
 - Experimentation: **Logged** interactions and provided new recommendations and hints.
 - User Feedback: **Rated** tools on ease of use and quality, and collected feedback through interviews.

Effectiveness of Multi-round Recommendations




Figure 1: Quantitative analysis on user study

Quantitative Analysis:

Unlike [Voyager2's manual exploration](#) and [LLM4Vis+'s time-consuming natural language queries](#). Experiments show that HAIChart [meets user analysis needs with lower interaction costs](#).

Effectiveness of Multi-round Recommendations




Figure 1: Qualitative analysis on user study

Qualitative Feedback:

User feedback shows that **HAChart is easy to learn and use**, helping users explore data and achieve high-quality visualizations efficiently.

THANK YOU & QUESTIONS?